Aquinas, AI, and the Pursuit of Learning

By Alex Stevens

My eight-year-old daughter recently invoked my fatherly duty to help her do research for a presentation about a figure from the Middle Ages; Thomas Aquinas, in her case. My wife and I bought a few books, checked one out from the library, and found some relevant material in reference works we already owned. Notably, we did not search for random articles or videos off the Internet or try using generative AI as a personal tutor; yet this was not enough to entirely avoid *AI slop*.

Among our stack of works on this historic figure, two stand in stark contrast – G.K. Chesterton's *St. Thomas Aquinas* and a cute looking, illustrated children's book. One, a well-respected ode to a man whose intellectual contributions can hardly be matched; a poetic sketch of a humble saint. The other, a simulacrum of a fun biography about Aquinas full of duplicate pages, dubious content, emoji bullet points, and enough six-fingered people to send Inigo Montoya into a frenzy. The tell-tale signs of AI-generated content were on full display in this children's book. These distortions of reality ended up being a rather ironic tribute to the brilliant mind that wrote disputations on truth and knowledge.

There was an ongoing joke while I was taking an Artificial Intelligence course in college – make sure your AI algorithm can't produce Skynet. James Cameron's *Terminator* has not left the popular imagination, but it is not the sole source for imagined future battles or dystopias. One of the struggles in *The Matrix*, to know what is real, makes an indefatigable T-800, or a T-1000 composed of mimetic polyalloy seem only mildly challenging by comparison. The rapid advance of generative AI has caused people to reconsider what would constitute science fiction (emphasis on fiction) in the next ten years.

It is that similarity with the challenges of *The Matrix* that, I think, causes anxiety about the future with respect to learning. There are real challenges in the pursuit of learning that are compounded by the wide-spread ability to quickly produce shoddy AI-generated content. It is easy to see how this evokes the question, “What is even real anymore?” or “Will people be able to separate fact from fiction?” While some will certainly be misled by it, this content does not fundamentally impair our God-given capacity to learn about His creation, to pursue the truth.

Eggs are Eggs

Even before generative AI became popular, learning without the benefit of a trusted instructor or instructor-selected materials could be intimidating. For many subjects, there are an overwhelming number of sources to choose from. Different authors can come from different schools of thought or emphasize a different set of facts, while some are just clearer than others. Not all works are equally trustworthy, either. Some have exaggerated or unsubstantiated claims, faulty logic, or even a valid argument based on a false premise. For example, how much credibility should be given to the claims of miracles after Aquinas’ death? Does the fact that something so extraordinary was only mentioned by one of my daughter’s sources reflect differently on the author who included it versus those who excluded it?

Another layer of complexity is context. People, events, and ideas do not generally appear *ex nihilo*. Knowing about King David’s time as a shepherd when he was young shows where he draws the imagery for the beloved Psalm 23\. Grasping the early trinitarian controversies requires knowing about the Arians, Nicenes, and debates about “*homoiousios*” vs “*homoousios*”. Understanding some differences between the Franciscan and Dominican orders colors Thomas’s family’s reaction to his choice to join the Dominicans.

The piles of books one needs to read to understand an argument and its historical significance can quickly become overwhelming. Yet we needn’t cry out, “What is truth?” or “Who will rescue me from this corpus of death?” Truth is knowable, even for non-academics. You can trust me, I've read people who have read Aquinas. As Chesterton wrote in his biography on Aquinas, "the philosophy of St. Thomas stands founded on the universal common conviction that eggs are eggs." This is as it must be for the Christian who believes that God has spoken to His people through His word, which was recorded for all time in the Holy Scriptures. God would not have given His word to His people and made them a pillar and buttress of the truth (1 Tim 3:15) if it was not knowable, even if all matters are not alike plain.

In the Psalms David compared God to a rock (Ps 18:2; Ps 62:6-7). He used his real knowledge of the world – rocks are hard, rocks are good for building – acquired through his senses and experience, to proclaim divine truth – the Lord is a fortress, a deliverer – acquired through his experience of salvation. Jesus taught about the kingdom of heaven through many comparisons (Matt 13:24–50; Mark 4:26–29; Luke 13:18–21). The objects (e.g., mustard seed) and experiences (e.g., finding a pearl of great value) he used for comparison would have been understandable to His hearers, even if they were blind to the truth of the kingdom. The apostle Paul opens his letter to the Romans with a bold declaration that creation itself makes it evident that there is a Creator, and, therefore, all are responsible for knowing it (Rom 1:18-20). So, even if some deny it, we can know that eggs are eggs and there is a good God who made them.

AI Slop Enters the Chat

While we can explore, learn, and know a great many things, that doesn't make all things equally easy to learn. Given the many challenges to learning that existed prior to generative AI, does the presence of AI slop make the challenge of pursuing truth insurmountable? Or, to rephrase the question, does the existence of content that is noticeably low quality, produced with presumably low effort, so completely hinder our ability, individually or collectively, to learn that it will cause a crisis? I don't think so.

Feed-driven social media, those platforms that are designed to keep you scrolling their feed (e.g., Facebook, X, YouTube, etc.), are probably where this content flows most freely, and I am not original in saying that it would do most people good to stay off it entirely. Given that many choose to stay on it, though, it would be to our great benefit to not treat it as a valid source of information. The effort to sift through what's found there will likely outweigh the benefit of searching there to begin with. After all, not all AI-generated content is noticeable instantly, and the remaining content is rarely known for its virtue or truthfulness.

Anecdotally, books are not even a guaranteed safe haven from AI slop. The relative ease to self-publish, while good for genuine writers, also facilitates these less desirable products. Even though retailers like Amazon make returns and refunds easy, encountering enough lackluster books will create issues of trust. Reviews, while not completely unhelpful, have their own trust problems as anyone who has bought a previously-unknown-brand item can likely attest. Like feed-driven social media, review-driven retailers might not be the right place to start a search when looking to learn.

While it is good for some people to know about looking for vanishing points from converging parallels or feeding visual noise through Fourier transformations to analyze images as Hany Farid demonstrates, this will not be accessible to everyone. The detection techniques are not likely to stay the same over time, either. Some are working on techniques that might be helpful for websites more broadly. One such proposal looks to create networks of sites that have been vouched to be human content. It doesn't hurt to hope that one of these solutions could become as ubiquitous as CAPTCHA systems in separating the wheat from the chaff. Adoption is always a struggle, though, and these solutions could find themselves with similar trust problems as reviews.

While it might be tempting to just disregard all articles, books, and videos published after 2022, when ChatGPT was released, this will not solve our learning problems. Time will go on, new topics will arise, advances will occur, and not all subjects will have the benefit of decades or centuries of publications.

Eventually, content from the AI era will need to be consumed, and that will be sooner rather than later for some fields, like medicine. It is doubtful, for example, that the broader population will advocate for doctors to not read any research on new strains of viruses simply because it was written in the AI-era.

Fortunately, AI slop is not necessarily introducing new problems so much as adding to old ones. From a learning perspective, the main problems with slop are that it is low quality and of dubious trustworthiness. These are real challenges, but they are not insurmountable. Further technology might play some role in overcoming it, but it seems more likely to me that cultivating learning practices whose first step is not "search the Internet" will fare better.

Truth, Trust, and Telos

For all learners, but especially for those who are self-directed, we should diligently seek the truth, think humbly in community, and remember our purpose for learning. The seriousness of the topic or level of interest might determine the rigor with which these practices are applied, but they can apply broadly, whether it is a hobby, career skill-up, or Bible study. Over time we can cultivate trust in different spheres – individuals, communities, publishers, etc. – that will help us avoid or mitigate the woes of the AI slop slough.

- *Diligently seeking the truth* is a commitment to live in and understand the reality God has given to us, not the one we would rather have. Asking questions, reading broadly, and validating what we learn (as much as we are able) are helpful ways to approach this.
- *Thinking humbly in community* is an acknowledgement that our own perspective can be limited and that there are many things we can learn from others. Soliciting recommendations and engaging others by articulating ideas and receiving feedback assume that we can learn from others. Humility, a shared goal or interest, and an appropriate level of openness will be the key to building trust.

- *Remembering our purpose for learning* is, hopefully, a freeing reminder that only God is infinite and omniscient. We do not have to go down the infinite rabbit holes we come across to say that we know, or know about, a topic or to be able to accomplish a particular task. Further, what is the “why” behind our “why”? For any pursuit, once going back enough levels, we should remember that our chief end is “to glorify God, and to enjoy him for ever.” (WSC 1). If the journey is at odds with that chief end, it is off track and needs to be righted.

These practices are not linear, mutually exclusive, all encompassing, or even equally applicable to all pursuits, but they are generally important to learning and building trust. Part of the problem of the pursuit of truth is lacking trust in the sources. Trust is predicated on dependability. Generally, the more often a person, group, or organization provides useful input (sources, feedback, ideas) that points us towards truth, the more we trust. AI slop is bad for the economy of trust not only when it contains false information, but especially if the recipients discover that the content is AI slop without a previous disclosure.

Given this relationship between trust and the pursuit of truth, it might feel like it is a chicken-and-egg problem. How can one learn without first having trust, and how can one validate a source if they don’t yet have knowledge? It is important to realize, though, that we have not lost our tie to reality and our God-given capacities. There are certain things we can verify through our senses, trial and error, and reason. Deriving everything from first principles would be quite a task, one that most are not up to; but we aren’t left to derive everything from first principles.

We started with a daughter looking to her parents for guidance. Children learn to trust their parents because they are (hopefully) consistently there to provide for them. We ought to be good stewards of that trust and teach them good things – truth about the world, the God who made the world, and how to live in it. As part of this, it is our privilege and duty to equip them for being in community and learning, for pursuing the true, the good, and the beautiful.

Lifelong learning is good, and I commend it to all, but there is an important caveat in this whole matter – there is a real need and place for experts, and, by definition, amateurs are not experts. My point is that one does not need total knowledge with absolute certainty on every point to say that they have knowledge. There are many endeavors we can undertake that do not require us to be experts, simply more knowledge than we currently have; ones that help us to better love God and neighbor, or be better stewards of what we've been given, or otherwise seeking to glorify God and enjoy Him. These endeavors are worth pursuing.

A Christian perspective should give us confidence in persevering in pursuit of learning, because it is not a fruitless endeavor. AI Slop is not an impenetrable barrier between us and God's creation; it is noise that can be overcome through diligence and community. Joining or forming communities that value the truth can be powerful for building trust, not because they do or don't use generative AI, but because they are tethered to reality. Welcoming requests for recommendations or questions that could have been a Google search might feel inconvenient at times, but it is more dignifying and constructive than the alternative, even when no one has all the answers.

Attentive Parenting in the AI Age

By Scott Hurst

Some time ago, I watched the promotional video for the Apple Vision Pro. Vision Pro is a spatial computer Apple plans to release next year. The Vision Pro goggles layer the world of our phones over the physical world, integrating the virtual world with the embodied world.

I had mixed reactions to the promo. I'm not qualified to critique the product so I can't say if it's good or bad on a technical level, but after the promo, I was both excited and troubled. My primary concerns were for my kids. As a teenager listening to music on my iPod Shuffle, I never thought I would need to shepherd my future kids through virtual reality. I was equipped for many situations as a parent, but not for a world with advanced robotics, A.I. and computers like Vision Pro. Those were for the movies, not concerns for discipling my kids.

Yet, that is the world we live in. It's the world my sons are growing up in. It's the world where they will need to learn what it means to be human and to follow Jesus. Part of my role as their dad is equipping them to love Jesus while living in that world.

Living well, at any time, requires paying attention. Paul makes this connection in Ephesians. "Pay careful attention, then, to how you walk—not as unwise people but as wise— making the most of the time, because the days are evil. So don't be foolish, but understand what the Lord's will is."

Attentiveness is essential for wisdom. Attentive to God, to the world and to ourselves. As a dad, I want my sons to be wise and Christ-like while living in the world of Vision Pro. For that to happen, I need to help them be attentive.

Be Attentive to “Messy” Experiences

I love big family meals, the ones where you need the oven, stovetop and grill to prepare. Every spot on the table is filled with food, and the signature meat sits as a succulent centerpiece. The kind of meal someone always captures for Instagram before eating. Those are my favorites… well they would be if I didn’t have to clean up.

What if we could have the meal without the mess?

Imagine sporting events and concerts without the smell of sweat and the feel of sticky beer-stained floors. Imagine the thrill of front-row seats and avoid public washrooms and crowded subway trains. Imagine a live stream so immersive that you get life without the mess.

I’ll admit, watching Toronto Raptors games on a living room-wide screen at home is enticing. But I’m cautious. Most digital experiences already take us away from the embodied. It’s rare to find any parent watching the school concert without their phone. We capture and edit those moments but struggle to remember the real experience. We got the video, but the sights, sounds, and smells of the room are gone because we were too busy working the camera to be present at the moment.

Our devices distance us from the moment and distract us from people. Standing in line at the grocery store or sitting on the bus without our phones sounds unbearable. Quietness, boredom, and meeting strangers seem to be burdens our phones liberate us from. We are constantly distracted and incredibly lonely.

We connect with others by eye contact. Eye contact can be super awkward and intimidating, but it also shows respect and attentiveness. It says, “I am listening, I want to hear you.” It is impossible to have our eyes on a screen and simultaneously make eye contact.

Products like Vision Pro, and Google Glass before it, attempt to solve this with a feature that lets someone see your eyes through the lens. But at best a computer over my eyes just splits my attention. Even now, simply having my phone in my pocket or on the table is a distraction. Why would strapping a computer on my head be better? Digital walls, no matter how transparent, are a constant threat against us seeing, knowing, and loving real people. Life without the mess isn't paradise.

In the gospels, we see Jesus in the mess, not above it. In one case he gives a blind man sight by spitting in his eyes\! The incarnation is when, as John tells us, the Word who is God became flesh and dwelt among us (John 1:1-18). He entered the mess. Jesus reveals a life of love and joy in a mucky world. Love led Jesus onto the muddy roads in Jerusalem.

In a massive meal, the mess is part of the experience. Every pot and pan on the stovetop and dish roasting in the oven fill the home with the aroma of the coming meal. Even the clean-up after is a family affair. With laughter and conversation drowning out the sound of scrubbing dishes and sweeping floors. The mess is part of the joy.

There is something human about unsanitized and messy experiences. We are meant to take in beauty, to be present in the beautiful moments of life. The best things in life are not always clean and not always recorded. The best moments are the ones with sore throats, a little mess, and a good friend. Real life has sticky floors and washroom lines. The relationships that matter and the experiences that make us are often messy, usually painful, but absolutely worth it.

Be Attentive to Your Moral Muscles

Moral muscles help discern good from evil and right from wrong. Training them strengthens our resolve to love God and love our neighbor. Advances in technology, though, often move faster than our moral training. In *Christianity and the Crisis of Cultures*, Joseph Ratzinger says, "The true and gravest danger of the present moment is precisely this imbalance between technological possibilities and moral energy." Technology moves so fast that moral formation chokes on its dust.

The speed and scope of new technology demand we develop our moral muscles. But what happens when the very technology we use weakens them?

Samuel James, notes how the environment of the internet, not just the content, forms us in harmful ways. “The habit of consuming content is a spiritual practice that has formative power in itself. We are being shaped into certain kinds of people by the kinds of things we linger over, including the things that disgust or grieve us.” Constant conflict is profitable for Facebook. Constant conflict is harmful for Christians.

Paul’s counsel to Timothy is helpful. “But reject foolish and ignorant disputes, because you know that they breed quarrels. The Lord’s servant must not quarrel, but must be gentle to everyone, able to teach, and patient, instructing his opponents with gentleness. Perhaps God will grant them repentance leading them to the knowledge of the truth” (2 Tim 2:24-26 c.f. Titus 3:9). Not every fight is worth fighting. For the sake of his soul and the sake of the gospel, Timothy had to stay out of the fray sometimes. Constant conflict breeds brawlers not wise Christians.

Our internet habits also don’t help. Attentiveness is the currency of moral formation but constant switch-tasking, skimming, and passive consumption sap our strength to be attentive. Let me give a personal example. I recently read an article about the war between Israel and Hamas, a conflict which requires careful thinking and deep reflection. Yet, the article was littered with ads for the best way to clean my phone. Not only did this break concentration, but the end of the article had three links to other articles on different topics. The message was clear: don’t linger over this article. Quickly jump to another. I gave up and watched TV.

If social media (where kids spend hours a day) makes our moral muscles flabby, how do we make them strong? Paul told Timothy that he had to pay attention to his life and teaching to make progress (1 Tim 4:16). This kind of attention requires training. Training the moral muscles in our family starts by saying no to some things. It means we don’t use an iPad only as a pacifier. It means explaining the “why?” behind discipline. It means providing resources and space for our sons and daughters to do the heart and mind work needed to form their moral convictions. It means giving them freedom to ask hard questions and doing our best to give them sound, thick and Biblical answers.

The moral formation of my sons starts with me. I need to turn away from my phone, to truly see the Godwardness of everything. If we cannot help ourselves from drinking the firehose of digital content, should we expect better from our kids?

As parents, we need to train our moral muscles well, so that we have a mature and biblical morality to entrust to our children. The Wisdom Literature in the Bible, which formed the core discipleship pathway from adolescence to maturity for a young person in Israel, serves this purpose. The aim of Wisdom Literature is summed up in Ecclesiastes 12:11, "The sayings of the wise are like cattle prods, and those from masters of collections are like firmly embedded nails. The sayings are given by one Shepherd." Work through these books with your kids. Have analog (face-to-face) conversations reflecting on their teaching. Make space in your family life for digital sabbaths, a time when all devices are off. Play board games, go for walks, or play football together. You cannot rush moral formation. Like all discipleship, it requires the work of reflection and unhurried stillness. You will not learn wisdom without making time to watch the ant (Pro 6:6).

Be Attentive to the Image of God in People

Putting computer chips into someone's body used to be science fiction. It may become normal before my sons hit puberty. Is the desire to become more like computers good for people? People want to become like computers. Is that good? Will it make us better humans?

My kids used to watch Ask the Storybots on Netflix. In each episode, kids ask questions about how stuff works and the Storybots go on learning adventures to find the answer. In one episode, the question is: How does the brain work? The Storybots answer, with a super catchy song, that the brain is a supercomputer.

There are similarities between human brains and computers, but this answer is problematic. Painting people in the likeness of machines blurs the line between human and computer to the point of non-existence. Our humanness is found in the differences between people and computers. Nicholas Carr makes this point in *The Shallows: What the Internet is Doing to our Brains*. “What makes us most human,” he says, “is what is least computable about us- the connections between our mind and our body, the experiences that shape our memory and our thinking, our capacity for emotion and empathy.” The computer comparison sacrifices too much. “The great danger we face as we become more intimately involved with our computers \- as we come to experience more of our lives through the disembodied symbols flickering across our screens \- is that we’ll begin to lose our humanness, to sacrifice the very qualities that separate us from machines.”

When the language of our inventions starts defining what it means to be human, we lose the givenness and Godwardness of humanity. Everything, even our identity gets created in our image and our value is decided by the same metrics I would use to judge my MacBook Air: usefulness, efficiency, and productivity. In this story, if I can’t use my legs to walk, I may become something less than human.

If we believe the highest good is about optimization then we will have no category for the value of weakness. Jesus’s call to take up our cross and daily die to ourselves will make no sense. When productivity reigns, people become tools to use and be used rather than persons to be seen, heard, and loved. Where does this leave people with disabilities, the elderly, or people deemed not smart enough or strong enough to be useful?

The praise of machine-quality usefulness is dehumanizing. The Bible tells a better story, one of humanity not made like machines but made in the image of God. People were created to bear his likeness and reign as his vice-regents over his creation. Our brains are not computers but reflections of God’s character and craftsmanship. Our worth is not in usefulness but is inherent in bearing God’s image.

The way forward is to be humbled by our unusefulness and learn habits of dependence on God. In *The Life We Are Looking For*, Andy Crouch says, “In an economy that evaluates and compensates us in impersonal terms, the most consequential members, the ones who matter the most for all our flourishing, are the ones whom Mammon does not consider useful. It is the “useless” who matter the most. Because if they are persons – if they are seen, known, welcomed, and given places of honor in our household – then all of us are set free from our usefulness.” When we stop worrying whether we and others are “useful” we are freed to see people as valuable. We learn to see people as God made them and we learn to love, serve and honour them.

The apostle Paul learned the value of frailty. He tells the Corinthian church about his humbling experience in 2 Corinthians 12:7-10

Therefore, so that I would not exalt myself, a thorn in the flesh was given to me, a messenger of Satan to torment me so that I would not exalt myself. Concerning this, I pleaded with the Lord three times that it would leave me. But he said to me, “My grace is sufficient for you, for my power is perfected in weakness.”

Therefore, I will most gladly boast all the more about my weaknesses, so that Christ’s power may reside in me. So I take pleasure in weaknesses, insults, hardships, persecutions, and in difficulties, for the sake of Christ. For when I am weak, then I am strong.

What some would call a liability, Paul learned to see as a gift. The thorn kept him humble and dependent. It reminded him that glory and power belong to God. Being human is not about transcending our limitations. Frailty opens our eyes to the beauty of being sustained by God through a living relationship with him.

When we teach our kids about being made in the image of God, we invite them into a better story. We invite them to find life-giving nourishment in union with Christ (John 15:1-11). We equip them not to sacrifice their humanity to computers, but to learn humble dependence on the God who made them. We equip them to love people, with all the frailty of humanity, as Jesus Christ has loved them.

Be Attentive to the Deceptive Gospel of Technology

We solve problems with technology. When we need what is underground, we make shovels. When we want to explore beyond the earth, we build spacecraft. With technology, we jump barriers and solve problems. This is often good (i.e. medicine), but our ambition can go too far.

Tony Reinke writes about the Gospel of Technology in *God, Technology and the Christian Life*. The Gospel of Technology is a worldview where our problem is not sin but nature. “According to the gospel of Technology, there is no fall of man, only impediments to the rise of man… Ultimately, whatever intrudes upon each person’s autonomy is the enemy, and the opposition can be defeated through innovation.” Technology plays the role of the conquering saviour creating a paradise for us.

The Gospel of Technology promises that every problem has a technological solution. It promises to bring about a perfect world through technological innovation. A perfect world at the tip of our fingers.

This desire was part of the human story long before Apollo 11 and the iPhone. The Tower of Babel reveals the ancient roots of the Gospel of Technology. “As people migrated from the east, they found a valley in the land of Shinar and settled there. They said to each other, “Come, let’s make oven-fired bricks.” (They used brick for stone and asphalt for mortar.) And they said, “Come, let’s build ourselves a city and a tower with its top in the sky. Let’s make a name for ourselves; otherwise, we will be scattered throughout the earth” (Gen 11:2-4). People at Babel used technology to supplant God and take his place over creation. Babel reveals what happens when technology wears messianic clothes.

Our ambition to create a paradise with our technology loses sight of one basic Biblical reality. Creation is under God’s curse: “The ground is cursed because of you. You will eat from it by means of painful labor all the days of your life" (Gen 3:17). Technology cannot solve God's curse.

When the Preacher in Ecclesiastes talks about his greatest achievements he shows us the personal paradise full of breathtaking palaces, magnificent gardens, and every luxury he ever wanted that he created. He created a type of Eden for himself and under his control. From the outside, it looked spectacular. On the inside, however, it was hollow. "When I considered all that I had accomplished and what I had labored to achieve, I found everything to be futile and a pursuit of the wind., There was nothing to be gained under the sun" (Ecc 2:11). His paradise was all smoke. Eden cannot be remade by technology. The hope of the new Creation lies elsewhere.

Nature's opposition is a consequence of human sin. That is why the hope of new creation is in the sin-conquering death and resurrection of Jesus Christ. God "condemned sin in the flesh by sending his own Son in the likeness of sinful flesh as a sin offering, in order that the law's requirement would be fulfilled in us who do not walk according to the flesh but according to the Spirit… For the creation eagerly waits with anticipation for God's sons to be revealed. For the creation was subjected to futility—not willingly, but because of him who subjected it—in the hope that the creation itself will also be set free from the bondage to decay into the glorious freedom of God's children. For we know that the whole creation has been groaning together with labor pains until now" (Rom 8:4-5, 19-22). Creation awaits a renewal by God's power alone because in the dying and rising of Jesus, God put sin to death.

Only the gospel comes with a promise that our bodies will be made imperishable, and all of creation made new by the power of God. Only God, through Christ crucified and risen, ends the curse (Rev 21-22). Hope in the face of sin and death isn't in technology. Hope is in Christ alone (1 Cor 15). To counter the gospel of technology, we need to keep in step with the gospel story.

My kids were born into a world of smartphones, live-streaming and Google Maps. This world assumes technology will eventually solve every inconvenience. We need sound Biblical Theology to challenge these assumptions and keep in step with the gospel. “If God is the center of your life, technology is a great gift. If technology is your savior, you’re lost.” Technology doesn’t have the shoulders to carry the savior’s cross. We need the savior who was born in a Bethlehem manger, not in the labs of Silicon Valley. To help set proper expectations of technology, help your kids see and savor Christ alone as Lord and Savior.

Be Attentive to Living by Faith

C. S. Lewis describes an approach to reality in *The Abolition of Man* that dominated the world of his day and the world he saw coming. In this approach, reality must be subdued to the wishes of men and “the solution is technique.” This approach sounds sinister but it strikes a chord with one of our deepest longings. The longing for control.

Our generation is not the first too long for control, but we express it a little differently. Chase Replogle makes a helpful connection between the use of sacrifices in ancient religions and our use of technology to gain control of our lives. “Ancient shamans searched for the means of appeasing divine forces, practices of worship, which might allow them to exert control over the chaotic experiences of weather, war, and fertility. Still today, it is the promise of technology, supercomputers with mass data and algorithms frantically scouring the universe for patterns that promise new possibilities of control and conquest.” Technology helps us control reality.

We can control our schedule, car, doorbell, Christmas lights, banking and more through a small device in our pockets. If we can design a product to control all of that, why not go bigger? What makes this desire for control a bad thing?

Craving control leads to destructive arrogance and fear. King Saul’s attempt to control God’s will made him paranoid and put murder in his eyes. Everyone became a threat to be snuffed out, even his family. The desire for control creates a life of instability. Today, we can create new identities and consume hours of content with minimal reflection but we’ve lost the stability necessary for peace. Today we are anxious, lonely and overwhelmed by the burden of creating our identity.

Maybe we aren't meant to be in control. Recognizing our limits is key to wisdom. Ecclesiastes 8:17 says, "I observed all the work of God and concluded that a person is unable to discover the work that is done under the sun. Even though a person labors hard to explore it, he cannot find it; even if a wise person claims to know it, he is unable to discover it." Some things will always be out of our control. Wisdom requires accepting and delighting in our limits.

The way forward is to press home the peace of resting in God by faith. Hebrews 11:1 says, "Now faith is the reality of what is hoped for, the proof of what is not seen." Faith is not surrendering our life to chance. Faith is a settled confidence in God. It is hope in the reality of his promise that is rooted in the integrity of his character. Instead of grasping for control with technology, faith courageously surrenders control to God.

Faith, planted in the stability of Christ Jesus, finds rest and peace. The bedrock of the gospel is better than the vapors of control. We can help our kids see this when we celebrate humble people of faith. Help them aspire to have the faith of the widow in your church instead of the money of Elon Musk. Shepherd them well by helping them choose faith over control.

When Paul walked around Athens "he was deeply distressed when he saw that the city was full of idols" (17:17). Parents today can sympathize. A quick walk through the digital world our children inhabit is deeply distressing.

We can learn from Paul. He didn't respond to this distress by checking out or giving up. Instead, he engaged. Luke, the author of Acts, connects Paul's distress to his engagement with one simple word: "So." "He was deeply distressed when he saw that the city was full of idols. So he reasoned in the synagogue with the Jews and with those who worshiped God, as well as in the marketplace every day with those who happened to be there" (17:17-18). Paul knew God was the Sovereign Lord. He knew God intends that people "might seek God, and perhaps they might reach out and find him, though he is not far from each one of us." He knew God "now commands all people everywhere to repent, because he has set a day when he is going to judge the world in righteousness by the man he has appointed. He has provided proof of this to everyone by raising him from the dead" (Acts 17:27, 31-32). Paul's faith in Christ caused him to be distressed by the idols of his time, but that same faith wouldn't let him check out.

His confidence in God, his love for Christ, and his conviction that the gospel is better than idols drove him to engage and call people to something better. I sympathize with parents who are overwhelmed by the technologies available to their kids. Parenting has always been hard and it isn't getting easier. But it isn't hopeless. Parents can engage and help their kids without fear and filled with hope because God is God, the gospel is true, and everything we need for life and godliness, God has already given us in Christ.

A Hollow Crown: AI and the Formation of Students

By Chris Sibben

The day after we showed our students *Forevergreen*, an Academy Award–nominated short film, something happened that I can't shake.

The film did what good art is meant to do: it quieted the room, even a room of K–8 students. It made them lean forward, their attention held by a sudden, rare gravity. Through the seamless marriage of image, music, and plot, it bypassed social static to reach their minds and hearts, holding them, if only for a moment, in awe.

Later, in conversation with seventh- and eighth-grade boys, I told them that the piece had been created by more than two hundred artists over five years. Their response came quickly, not as flippant dismissal but as baffled practicality: *Why?* One student said what the others were thinking: "AI could have done it in five minutes."

In those two small moments, the educational dilemma of our age revealed itself. Students have acquired more than a new tool. The technology is quietly remaking what they believe human work is, and what they believe a human being is for.

In the weeks since, I've heard a companion phrase everywhere. Show a new image, a new paragraph, or a new idea—something that once would have invited curiosity or debate—and someone will shrug: "That's AI." Not as a moral alarm, but as a conversational stop sign. A way to end the encounter before it begins.

Before going further, the tension deserves to be named.

AI is a threat to meaning. It is also a tool for redeeming time. Teachers use it to clear administrative undergrowth that crowds out actual teaching. Students use it to organize research, test an argument's weak points, and move past the blank page into the harder work of thinking. These are not trivial gains. Time is finite. Attention is precious. If a machine can shoulder the scaffolding, more of both can go toward what matters.

The question, then, is not whether to use these tools. If we want to understand the role of generative AI in education, we must first understand the ends these tools serve. And that, in turn, forces a deeper, more uncomfortable question: What is education actually for? We must answer this not in the language of outputs and metrics, but in an older, more demanding sense: What kind of person is a school trying to form, and what practices form that person?

To answer that, it helps to return to something older than our panic: the nature of signs.

Augustine of Hippo makes a deceptively simple distinction. There are things, and there are signs. A thing is what it is. A sign points beyond itself. Words are signs. So are gestures. So are sacraments. In a classroom, student work functions the same way. An essay is not valuable as ink on a page; it is a sign that points to a student's mind: his grasp of an argument, his habits of attention, his ability to name what is true and discard what is not. A lab report is not the point. It is a sign of whether a student can submit to reality, whether he can observe, revise, and be corrected by what is. The artifact is never the telos. It is evidence that something has happened in a person.

For a long time, that semiotic economy held because the sign was tethered to time. Work took time. Time left traces. Those traces could be read, imperfectly but meaningfully, by a teacher and a community. Even a mediocre essay carried a certain dignity. It bore the marks of limitation and struggle—the marks of a human being learning.

Generative AI breaks that tether. It manufactures signs without the slow formation those signs once implied. It delivers the look and tone of thinking without the ordeal of thinking. The change is more than the fact that the tool can produce an answer. It can produce a trace, and it can do so at scale. The very kind of object teachers have historically read as evidence of becoming is now cheap to generate.

Signs have always been detachable from their makers. A sentence can survive its author. It can be quoted, recopied, and recontextualized. The power of a sign is that it can travel, repeated in contexts the original speaker never intended. A signature can be reproduced. A text can be lifted out of its life. If a sign could not be repeated, if it were bound to one private moment of presence, it would not function as a sign at all. Meaning depends, in part, on this portability. Language works because it can outlive us.

So what, then, is new?

Not absence. Not detachment. Not even imitation.

What is new is the industrialization of detachment. Signs are produced at conversational speed, tailored to the user, and clothed in the social costume of personhood. The machine does not publish; it answers back. It adjusts. It apologizes. It seems to take responsibility while having no continuous self to bear the costs of its own words. It produces the kind of speech that, in ordinary human life, is bound up with vulnerability and obligation, yet it is structurally incapable of either.

Put differently, what is a permanent feature of signification—repeatability without the guarantee of presence—AI turns into an atmosphere. The classroom does not contain detached signs; it is saturated with them. And once detached signs become normal, students become unsure which, if any, sign can be trusted.

They may respond in two predictable ways: contempt and suspicion. One drains the work of dignity. The other drains the world of wonder.

This is not entirely new. Even Shakespeare saw what happens when the sign floats free of the substance. In *Richard II*, the deposed king stares into the nature of authority and discovers, too late, that he has confused the symbol with the thing:

For within the hollow crown

That rounds the mortal temples of a king

Keeps Death his court, and there the antic sits,

Scoffing his state and grinning at his pomp.

The crown is hollow not because it is fake, but because Richard wore the signs of sovereignty without the inner formation, character—even virtue—that would have made them durable. When the ceremony and pageantry are stripped away, nothing remains that can stand.

That is the risk we are running in education. The danger is not that students will use AI to cheat, though they will, but that we will gradually accept a world in which the signs of learning circulate without the formation that once made them trustworthy. Essays without writers. Voices without owners. Credentials without the slow ordeal of becoming.

The hollow crown is not the AI-generated paper. It is the graduate whose diploma no longer points to anything real.

For most of history, personal effort was the visible cost of meaning. A poem mattered because you could feel a person behind it: spontaneous insights, false starts, and long struggle, the refusal to let a line go until form and function united. The words were a kind of signature, more than ink: the trace of a life and a way of seeing. A painting mattered for the same reason. You could sense the time inside it, the hours and disciplines that leave their mark. Even when you could not name the technique, you could recognize the ordeal of seeing and making: an artist submitting to reality or rebelling against it, accepting limits or rejecting them, and then inviting others to see what he sees.

Schooling rests on the same premise. An essay is not valuable because it fills a page. It is valuable because it is a record of a mind learning to think, a heart wrestling with the other, a person who is real. The artifact is a sign. It gestures beyond itself.

But here is what we rarely say plainly: the crown being earned is not a credential. It is a self.

A student who has wrestled with a hard text, revised an argument under pressure, and failed and tried again is more than informed. He is more solid. He has learned, however imperfectly, to be answerable to reality rather than expressive of preference.

That solidity is the telos of education. Not information transfer. Not skill acquisition. Not even "critical thinking" in the thin sense the phrase has come to mean. The end of education, rightly understood, is the formation of a person who can stand somewhere: someone with a voice, convictions, and a relationship to truth that has been tested and is therefore owned.

Generative AI does not threaten that telos by being wrong. It threatens it by being easy. It offers the appearance of having thought without the ordeal of thinking. And the ordeal is not an obstacle to formation; it is the mechanism.

"AI could have done it in five minutes" is not finally about speed. It is a moral revaluation. It assumes that what matters is the output, not the ordeal; the image, not the seeing; the product, not the person becoming capable of making it.

Live in that assumption long enough and time itself begins to look like a mistake.

When time becomes suspect, education becomes vulnerable at its root. Schools have always insisted that the young endure what is not immediately gratifying: the resistance of a hard book, the humility of a wrong answer, the slow dignity of revision. They teach, often against every adolescent instinct, that truth is not produced by desire. It is discovered through attention, patience, question, critique, and resolution—a continuous, embodied encounter with what is real.

The AI environment trains the opposite instincts. It offers competence without apprenticeship. Fluency without understanding. It makes it possible to generate signs without the patience required to become the kind of person those signs once implied. A student who internalizes that pattern does not become lazier; he becomes less formed, less present, less able to bear the weight of difficulty without reaching for a prompt.

This is where C. S. Lewis offers an image that is difficult to shake. In *The Great Divorce*, the souls who arrive from hell are not monsters. They are shades—thin, translucent, barely there. They have spent so long avoiding the weight of reality, so long choosing the frictionless comfort of their own preferences over the hard country of what is, that they have become insubstantial. They cannot bear the grass of heaven; it is too solid for them.

The danger AI poses to students is not primarily moral in the narrow sense. It is ontological. A student who consistently outsources the struggle, who prompts rather than thinks, generates rather than wrestles, or accepts fluency as a substitute for understanding, is choosing the path of the shade. He is becoming less solid, less capable of conviction, less present to himself and to others.

And the tragedy is that he may not notice. The signs will still be there. The essays will be produced. The grades will be earned. The diploma will be conferred. The hollow crown will be placed on a head that was never formed to wear it.

This is why the question of when and how to use AI is not a policy question. It is a semiotic and pedagogical one. If the artifact no longer reliably signifies formation, schools must move closer to forms of learning that cannot be reduced to a generated trace.

That shift is practical. It looks like more in-class writing and problem-solving where the mind is visible in real time. It looks like drafts submitted with commentary on what changed and why. It looks like oral defenses in which a student must explain the choices behind his claims. It looks like lab notebooks that record failure, revision, and the stubborn confrontation with what is. It looks like teachers who assess not only a product but a trail.

The honest answer to when AI is useful for students is only after the productive struggle — when it serves their formation rather than replacing it. AI is useful when it clears away the mechanical after the student already understands the mechanics; when it helps a student see the gaps in her own argument rather than filling them for her; and when it accelerates the path to deep work because the student has already done the hard labor of finding her way.

But that answer only makes sense if we know what the hard work is for.

A school must be able to say, without embarrassment, that its purpose is not to produce graduates minimally trained for the economic and political order. Its purpose is to form people who can bear the weight of reality: people who can read with care, speak with honesty, revise under correction, and stand behind their words.

That is not a mystical claim. It is physical. Formation shows up in flesh and habit. It shows up in whether a student can sit still with a difficult paragraph without reaching for escape. It shows up in whether he can name what he believes and why. It shows up in whether he can endure the small humiliations of being wrong without collapsing into cynicism or defensiveness. It shows up in whether he can apologize without theatrics and repair without resentment. These are not “soft skills.” They are the contours of a person who can be trusted.

When a school succeeds in forming that kind of person, it produces something AI cannot replicate: a voice. Not a tone, not a brand, not a borrowed style. A voice in the deeper sense—the sense in which a person has something to say because he has become someone. His words carry weight because they point back to a self shaped by attention, constraint, failure, and love of what is true.

A student formed that way can use AI freely and wisely. He can delegate the scaffolding without surrendering the substance. He can move faster when speed is warranted, and slow down when slowness is the price of depth. He can use the tool without becoming its product.

Forevergreen gave my students a brief taste of that older gravity. Beauty felt costly, and therefore meaningful. The film reminded them, before they had language for it, that labor is not inefficient. It is dignified. The making of a thing is part of its meaning.

Then the room brought that gravity into contact with the new regime. Five years. Two hundred artists. *Why?*

Why read slowly? Why memorize? Why draft? Why become something when you can generate the sign of having become it?

The unsettling part wasn't the question. It was how normal it felt.

That is the mark of a regime change: not when a technology appears, but when its values start to feel like common sense. When time feels like waste. When effort feels like error. When simulation feels sufficient.

Richard's hollow crown is a warning, not a metaphor for despair. He lost his kingdom because he mistook the ceremony for the substance. We are in danger of making the same mistake, accepting the signs of education while quietly abandoning the formation that once gave them weight.

If schools forget that, they will still have essays. They will still have grades. They will still have endless, fluent signs. But those signs will no longer reliably point to a person. And once that link breaks, we will not only be unsure what student work means; we will begin to lose confidence in what words, promises, and judgments mean at all.

AI is Coming for Online Pastoral Education

By Cameron Shaffer

Artificial intelligence is going to make online theological & pastoral education obsolete by 2030\. That's a bold claim (feel free to remind me that I got this prediction wrong\!), but likely to come to pass. Online theological education as it is currently deployed may continue to exist by that year, but it will be hopelessly and clearly dated and on its way out.

There are currently two kinds of online theological education. The first developed to provide pastoral education where it was previously *absent*. This approach is aimed primarily at the developing world or countries with infant churches that lacked their own institutional resources. The second kind of online theological education developed as an *alternative* option for constituencies that formerly attended classes in-person or residentially. Both kinds, but especially the alternative model, are going to find themselves out competed by AI education.

I'm going to provide a description of current online pastoral education, a prediction for the effects of AI upon it, and prescriptions for how to adapt for the sake of the church.

What Is Online Theological Education?

Online pastoral education required the right technological conditions to emerge, most notably the rise of broadband internet. What we call "online education" typically means pre-recorded lectures streamed over the internet, sometimes augmented by live video classes, the best versions of which provide some kind of digital interaction between students and teachers. In many cases, however, it amounts to self-guided reading, online quizzes, and some form of grading, whether by a professor, a TA, or increasingly, a computer.

When it first appeared, online theological education had to justify itself as a credible alternative to the traditional in-person model. Within higher education this wasn't unique to seminaries, but the arguments made in this context were especially weighty. Churches had the appropriate intuition that pastoral education should be more than informational transfer and include a personal dimension. So when online models emerged, they had to explain how they weren't merely digitized content delivery, but a valid path to real pastoral preparation.

There are generally three arguments in support of online education as an *alternative* model. First, it offered access. Students who were unable to attend seminary due to job constraints, family commitments, financial limitations, or geographic isolation could now connect with high-quality faculty, gain institutional credentials, and benefit from curated curricula. The institution's reach now extended far beyond its physical footprint to recruit those who earlier would not have been able to attend; students no longer had to bear the expense of moving with all the attendant costs (social, ministerial, contextual, familial) associated with a relocation. And in a competitive educational marketplace, whichever schools move online first or best are able to more effectively recruit students who wouldn't otherwise consider them. Hesitation seemed like acquiesce to obsolescence.

Second, online education is cheaper. The more students were online, the less overhead (dorms, class rooms, cafeterias, landscaping, etc.) was needed. Class sizes could expand. Fewer full-time faculty were needed, and lower-paid adjuncts or administrators could facilitate the courses. As costs dropped, the appeal to price-sensitive students grew. Affordability opened the doors wider, which made theological education more accessible and an appealing alternative to the expensive, inconvenient traditional model.

But third—and most important for online *pastoral* education—students could stay rooted in their home churches and cultures. Churches no longer had to send their emerging leaders away, risking the loss of those future leaders altogether. And students didn't have to uproot their families or abandon their communities in order to prepare for ministry, but instead could be trained in the specific context where they were called to ministry in the first place.

This was the heart of the online pastoral education argument: the very technology that enabled distance learning also allowed for a more deeply embodied connection to the student's local church and context. Good online theological programs recognized that formation had to remain personal. Some schools added features to help address this, such as professor Q\&As, virtual office hours, mandatory online discussions in forums, and occasional cohort gatherings in-person or online. The best versions required a local mentor, often a pastor, to walk alongside the student.

Still, at its core, online education is a tool for transferring information. And while the informational transfer aspect of seminary education is frequently denigrated, it remains essential for pastoral preparation: pastors need to know the right stuff. Content delivery is essential for ministerial preparation, and has always been a central feature for pastoral training.

The difference is that the context of that transfer shifted from an embodied, relational classroom to a digitized avatar, often in a depersonalized format. Online education needs to supplement this deficit by being more personal and more in-person, and often does so, just typically not through the school itself, but rather through the student's church and mentors. Online education can still function well for pastoral education, especially in the *absent* model, because it curates information within the theological tradition of the school and student, provides assessment and credentialing by conferring degrees, and allows for students to apply what they are learning in their context.

But here's the problem: once you've moved online, you're no longer just competing with the seminary down the road. You're competing with *everyone*. Students can now compare your faculty, curriculum, and format with those of every other institution in the world. When all things are equal—when access is universal and the platform is the internet—the edge disappears.

The online pitch is: Our education provides good information cheap; you can stay home (convenience\!) and formation will occur in your church where you serve.

The AI pitch will be: Our education provides even better information more cheaply; you *still* get to stay home and formation will occur in your church where you serve.

AI Changes the Game

AI will outperform traditional online models because it isn't just distributing information: it's tailoring it. It will personalize content, adapt in real time to a student's learning style, provide instant feedback, and offer a level of interactivity and contextual relevance that prerecorded lectures and online forums cannot approach. It will be cheaper. It will be more scalable. And paradoxically, it will feel more *personal*, not less.

We're not there yet. But we're not far.

\[It\] means we could go in a relatively short span of time — a year or possibly less — from A.I. systems that look not that different from today's A.I. systems to what you can call superintelligence, fully autonomous A.I. systems that **are better than the best humans at everything**. In 'AI 2027,' the scenario depicts that happening over the course of the next two years, 2027-28 (emphasis added).

This is Daniel Kokotajlo, Executive Director of the AI Futures Project, this past week in the *New York Times*. The next 18 months will likely witness an exponential growth in AI capability that will be absolutely mind-boggling.

The difference between current AI and what's coming is greater than the difference between the first rotary phone and the latest iPhone. The next generation of AI will be able to read everything that's ever been written, watch every recorded lecture, and listen to every interview. It will absorb every new book, every journal article, every publication as it's released. It will not forget information it comes across. It will synthesize all of it instantly and deliver the information more thoroughly and accurately than any human professor ever could. You don't have to believe that AI will possess actual intelligence, creativity, consciousness, qualia, sapience, or personhood to see the challenge here. AI, for all practical purposes, will be a self-possessed infinite library capable of guiding students to its riches.

And AI will no longer just wait for prompts through a text interface. It will be able to converse in real time, and offer guidance proactively, like a theological tour guide. It won't just generate text, it will appear and speak. Deepfakes and digital avatars—already impressive—will become indistinguishable from real faces and voices.

The AI professor won't just present content. It will be compelling, persuasive, and even charismatic — and more so than humans. Video classes powered by AI will feel more engaging and personal than any Zoom call or pre-recorded lecture ever has. Its ability to adapt to the student, on the fly, will outpace every digital classroom we've built so far. And even if governments or regulatory bodies are able to require things like AI signatures on these presentations, they will not dissuade students any more than a watermark or institutional logo currently does. Students will be studying (or whatever you prefer to call it) with the AI precisely because it is AI.

And the backend won't escape either. AI will collapse the administrative load that currently bloats online education—grading, paperwork, and bureaucratic processing will all be automated and immediate.

All of this makes AI-based education cheaper than human-based online education. It also changes the competitive landscape. Right now, online schools have to compete globally. Once AI fully enters the scene, they'll have to compete universally—not just against the current faculty in other schools, but against all the voices of history. St. Augustine and Martin Luther. Athanasius and the Cappadocians. C. S. Lewis and Karl Barth.

AI will not only quote them, it will have “read” their complete works and be able to interact with them more fluently than any of us, all the while looking and sounding like them. Your online professor might know their material brilliantly, but AI will allow you to learn the theology of the *Summa* “directly” from Thomas Aquinas or how to preach from Billy Graham.

The difference between AI-mediated and directly learning from past figures will seem immaterial and theoretical to students. Protest the inhumanity of it all, bemoan the loss of the real; it won’t matter. All of the reasons that students opt for online education (access to information and institutional credentialing, affordability, staying rooted) and why schools promote it (increased student base, lower costs) not only remain in AI-education, but are strengthened.

Even if the most bullish predictions of AI’s capabilities in five years are overblown, the pressure from AI upon online education is already intense and only going to increase. There’s no way the current approach to online education can keep up. This is not just another shift within the same model. It is a transition to a new era. To use an analogy from an earlier era of the internet, online theological education is not Barnes & Noble taking over the local independent bookstore. It is Borders being overwhelmed by Amazon.

What’s likely to happen is that schools and ministry platforms will subscribe to AI packages, trained within and tailored to their theological traditions. Much like subscriptions to streaming platforms (Amazon/Audible, Spotify) and digital library service bundles (ATLA, EBSCO, Logos) AI will not circumvent copyright, but become intertwined with it. These AI engines will become a new online faculty, curating and delivering the education, answering questions, guiding reading, and walking with students. A subscription package is going to be a lot cheaper, and more effective, than a fully staffed online department.

Online education as we understand it will not be destroyed by AI. It will be absorbed, replaced, and quietly left behind.

What's to Be Done?

But not all hope is lost for schools that invested heavily in online education. Barnes & Noble, once on the verge of collapse, is now in the midst of a quiet renaissance, not by outcompeting Amazon, but by going local, reinvesting in in-store experiences, and cultivating what Amazon could never replicate. The same possibility is open to theological schools who have embraced online education: adapt not by resisting AI, but by leaning into what it cannot do. Schools that prepare wisely for the AI revolution can pivot successfully. More than that, the church needs them to.

So how should they prepare?

First, online seminary programs shouldn't treat AI as the enemy. AI replacing online courses is not a loss—it's a refinement. Aside from the issue of pastoral formation already being handled in local church settings, all the other benefits of online education are provided more efficiently, simply, and cheaply by AI. Information transfer has always been a key part of ministerial education, and a model that facilitated digitized information transfer being displaced by one that does it better is not a philosophical disaster. It's progress. If a school has already crossed the digital Rubicon, there's no point trying to paddle backward by holding the line at Zoom. AI will do what online courses were designed to do, only faster, more thoroughly, more accessibly, and at a lower cost. The wise move for online classrooms is not to resist AI, but to offload content delivery to it with intentionality, and then reinvest in what AI cannot do.

Many online schools and faculty will protest that, even though it may be done by livestream, there is still a face-to-face interaction between human faculty and students, which will remain indispensable. Interaction between *people* is crucial for formation, not just as pastors, but as humans. Megan Fritts, professor of philosophy at the University of Arkansas Little Rock, in a beautiful recent essay made the case that there remains a formative value for a humanities education because "the very essence of humanity entails something that, in fact, \[AI\] will not do at all." Invoking Ludwig Wittgenstein and Alasdair MacIntyre, she argues "Human linguistic expression arises *from* the human form of life and also *shapes* the human form of life. There is arguably nothing more distinctly human about a life than that it is filled to the brim with language." Language and conversation and interpersonal interaction are necessary for human formation, and currently ChatGPT cannot provide that. The online interactions between faculty and students remain formative, still. But where I fear Fritts' argument falls short is that AI will soon be able to speak *like a human to humans "better" than humans*.

I did my pastoral education in a traditional, in-person institution, but my doctoral work is a combination of self-driven study and online interactions with my advisors and learning cohort. Those online, Zoom-mediated, face-to-face conversations are invaluable for my education, but the reason they have, and will continue to hold, real value is because I have met all of my digital interlocutors in the flesh.

Second, institutions must re-center their programs on personal formation. Any digital interactions between human faculty and students need to be anchored in some degree of in-person connection: hybrid courses, week-long in-person intensives, annual retreats, visits from the faculty to students in the *absent* model. That means giving serious attention to mentoring, local church involvement, and student peer cohorts. The strongest online programs already understand this and supplemented their coursework with meaningful in-personal components. Those features must now move from supplement to center. The marketing shift should not be "We are better than AI" but "Here is our AI-powered curriculum, and here is how we ensure you have embodied, personal formation alongside it."

The school's value-add becomes less about the content (though that institutional and theological-traditional curation for the sake of credentialing will remain important) and more about the community and the context in which that content is lived and ministerially brought to bear. This will be especially true since students will wonder why they need to pay tuition if they can just purchase the AI-subscription package on their own anyways. Digital faces, AI or otherwise, in pastoral education will all be the same to students unless there is an embodied connection between teacher and learner.

Third, online schools must aggressively localize their student experience. That means forming cohorts not just around shared classes, but around shared geography and ecclesial life. Churches that welcomed online sermons and video classes are likely to welcome AI-driven tools as well. But they will need trusted pastors and theologians to walk with students through them if they are be fruitful in the life of the church. Schools can help by building networks of mentor-pastors and regional facilitators. The best online programs will not be the ones that resist AI, but the ones that embed it within a larger structure of discernment, discipleship, and pastoral accountability. Sermons should be delivered with a people, known and loved, in mind. AI may be able to provide the best rhetorical guidance, but congregations are known personally. Mentors in the local church are essential for helping students move from speakers to embodied and embedded *preachers*.

In short, theological schools must stop thinking of themselves as the source of all education and begin positioning themselves as curators of formation. AI will soon become the primary tool for delivering theological content for the everyman in the pew or desk. That is not where the battle lies. The battle will be in who provides the context, the community, and the wisdom to help students discern what to do with what they're learning. That cannot be automated or downloaded. That must be lived, face to face, in the body of Christ.

What About Traditional Seminaries?

In their best versions, online pastoral education prioritizes formation within the student's current ministry context; in-person education treats preparation itself as the student's vocation. That means that all the embodied aspects of online education (learning cohorts, access to pastors and mentors, connection to the local church) need to remain present in and be stressed by traditional schools. Similarly, traditional schools are not exempt from the challenges ahead. In fact, some of those challenges are about to intensify.

Residential theological education will soon be, comparatively, dramatically more expensive, even when underwritten by generous donors or institutional endowments. And it will be costlier in terms of time and energy. As AI becomes capable of delivering more accurate, responsive, and personalized instruction, schools will have to work harder to explain why students should leave their homes and ministries to pay significantly more for something AI can provide more efficiently.

That means traditional schools must make a clear and compelling case. They cannot assume that their value is self-evident, nor can they rest on nostalgia or prestige. They will need to show why learning in-person, in a classroom, at a slower pace, with embodied faculty and peers, offers something essential that AI cannot replicate and is worth still receiving instead of what AI can provide. Fritts argument from the formative nature of language in community is one.

But another way to begin making that case is by drawing from what Pierre Bourdieu called *habitus* and what James K. A. Smith has framed as the formative power of liturgy. The sets of dispositions, habits, and ways of being that we absorb through repeated practices within a community shape us. Ministerial formation is not only about what we learn, but about *how* we learn it. It is about the posture we take toward the truth, the community in which we receive it, and the pace at which we are shaped by it. In-person education allows for a kind of deep formation that comes not just through content but through environment, repetition, and the embodied presence of others; by making that the vocation and rhythm of life for the student. In other words, it is not just how information is received, but how it is *given*, that matters. That interaction is personally formative, and especially in the context of theological study, is spiritually formative.

In this model, the classroom is not merely a delivery mechanism for information. It is a space of communal formation. It is where the future minister learns to attend to God’s word with others, to hear truth not only in words but in tone, expression, and shared life with a fallible professor whose quirks and shortcomings cannot be adjusted at a whim. It is a slower process and less efficient one, but it is precisely in that slowness and inefficiency that something vital is formed.

Personal, embodied teaching of the gospel prepares students for personal, embodied ministry. The challenge for residential schools in the age of AI will be to show that their slower, more expensive model is not a relic of the past but a deep investment in the future of the church. They will need to articulate not just what they teach, but how their way of teaching is part of the formation they offer. That will require clarity, humility, and a renewed commitment to the kind of formation that only happens when truth is shared in person, among the people of God.

In all these models, the new reality is that the theological institution is no longer the primary dispenser of content, but still has an invaluable role in pastoral formation as the guardian of tradition, mentor for maturity, and the cultivator of character. This is work no AI can do. And it is exactly the work the church most desperately needs seminaries to do for the sake of Christ's sheep, whether its pastors are formed via online-facilitated or traditional education. In the coming days, an embodied ministry of the gospel is what will distinguish the church from an AI-dominated world.

Pandora, AI Girlfriends, and "Reborn" Babies

By Nadya Williams

Once upon a time, very long ago, a man lived alone. Perhaps he wished for a companion, perhaps not. We really don't know. What we do know is that for a set of very complicated reasons, when the well-being of mortals became a pawn in the power struggles of the pagan gods, the gods created a wife for him. She was, in some ways, the perfect woman, a collaborative project of the gods, who each endowed her with a gift, resulting in her name: Pandora—"all the gifts" or "the gifts of all."

Except, she was not a woman; she merely looked like one. She was, rather, an automaton—or a robot, if you prefer. In her book Gods and Robots: Myths, Machines, and Ancient Dreams of Technology, Adrienne Mayor looks at both literary and artistic portrayals of Pandora, and finds a consensus: Pandora, the product of the workshop of the craftsman god Hephaestus, was seen in the Greek imagination as an artificial being, rather than real flesh-and-blood woman. As a machine, she was designed for a purpose, and that purpose was nothing good.

Chances are, you have heard about Pandora and the ills she caused—plagues upon the world that she let loose out of the container that was her twisted dowry from the gods. These ills are a fitting reminder of the warnings that already ancient pagan writers had in mind when dreaming of technology far beyond anything possible in their lifetime. We can create artificial beings that may look like us—sort of, on the outside. But these artificial beings deceive, distort, and destroy those with whom they interact. Mayor explains:

Pandora calls to mind another myth about a cunning artifice that was a dangerous gift—the Trojan Horse. Some versions of the story of the Trojan Horse, built by the Greeks and presented to the Trojans as a ruse of war, suggest that it was sometimes imagined as an animated statue with articulated joints and eyes that moved realistically…As a being that was made, not born, Pandora is unnatural. A replicant with no past, Pandora is unaware of her origins and her purpose on earth.

Robots and AI creatures cannot die, yet they cannot live either. They exist outside of time, for they have no real memory, no ability to feel and experience events the way real flesh-and-blood people do. Their realistic appearance is but a weapon of deception and destruction.

Epimetheus, the one who received the cursed gift of Pandora from the gods, had no idea what he let into his home. He was fooled simply because he was foolish—one whose name literally means "after-thinker."

As more Pandoras proliferate around us, however, not everyone is fooled by accident. Some today choose the lie, preferring it to the truth. The consequences are no less destructive than those imagined in antiquity.

Recently, a report on the rising trend of AI girlfriends made rounds. What is an AI girlfriend? She is anything you wish her to be. "Eva AI invites users to create their dream companion, while Dream Girlfriend promises a girl that exceeds your wildest desires. The app Intimate even offers hyper-realistic voice calls with your virtual partner." In a world of customer-is-always-right, an interested consumer can now design the perfect woman for himself. There's just one catch: she won't be real. She will be virtual. But close enough—she will sound real, at least. Critics have rightly noted the problem with distorting men's ideals of relationships through such a process. But there is more.

This fantasy world of creating relationships outside the realm of real life but mimicking reality to a grotesque degree is not limited merely to romance. Four years ago, The New Yorker profiled "Women Who Mother Lifelike Baby Dolls." Termed "reborns," the dolls look realistic to an extreme degree, mimicking real-sized infants. Only up-close does one notice that they are not real flesh-and-blood. They are not dead, but they are not living either. Still, for those who collect them, they seem deeply real:

A woman named Marilyn told Harris that “the dolls here really are like part of the family.” Another, named Laurel, said, “Holding the baby literally plays with your mind, it’s so real.” And Kym, a stepmother of five, wanted to have a baby who would be only her own; when she ordered a reborn doll online she “fell in love with it, because it looked and felt like a real baby, the baby I’d never had.”

What might be the problem with these relationships—AI girlfriends or reborn babies? Should we be quick to condemn these phenomena, considering that such practices are not, technically, harming anyone? I contend that yes, we should, for these practices stem from a worldview that rejects the preciousness of human persons and sees created beings—virtual girlfriends or plastic babies—as equal to real people. In fact, individuals who choose these replacements for live people usually do so in lieu of such relationships in real life.

We know that what we love shapes us. Loving replacements, replicas of reality, shapes us to embrace the lie as if it were truth, only leading us further from a life lived in real family, real community. Such is not a life of flourishing.

Yet in some ways, the temptation of these “relationships” is understandable. The ads for AI girlfriends play on this overtly: as they say, here is the dream girlfriend who will never disagree, never criticize, never disappoint. The subtext, of course, is: any other girlfriend will. Or, in the case of a baby, here is an infant who will never wake you up at night, will never spit up or get sick and make you miss work, will never just inexplicably cry for hours on end as your nerves fray from exhaustion and pity and guilt, as you pace the room trying to comfort this crying infant, sobbing yourself and thinking that you will never ever sleep again. But what do we miss when we settle for such “relationships”? We miss humanity, but we also, in the process, miss God, whose image we get to encounter in every imperfect human being whose we love in this life.

It is tempting to blame our twenty-first century society for such twisted desires to play God and create people—fake loved ones—in one’s own image and for one’s own sole pleasure, all while disassociating ourselves from the difficult and demanding people with whom we could instead lovingly live our lives. It is no accident, however, that stories like that of Pandora survive already from antiquity as cautionary tales, rather than celebrations of technological innovations and automata saving the day.

Dreams of automata stem from dreams of control. Should we really desire people to be only what we want them to be, rather than seeing them as God's image-bearers in their own right? How bad would the results be if we let someone continue down the road of AI girlfriends or relationships with automata? Here is one modern cautionary tale by way of an answer.

In 1884, the young French writer Rachilde (her pen name) published one of the most scandalous novels of the day, Monsieur Vénus. In this book, an overly bored aristocrat, Raoule, treats an impoverished florist, Jacques Silvert, as the object of her pleasure, marrying him but molding him further and further into her own vision of desire. At last, when he is killed in a duel—which Raoule had orchestrated as well—she orders a lifelike automaton to be made as his replica. Not going so far as to embalm Silvert's body, she orders nevertheless that his hair, eyelashes, and teeth be extracted from his corpse and installed in the automaton. It is with this robot replica of Silvert that Raoule continues to be amorously involved at the close of the novel. Having killed the human version of the man with whom she had toyed so cruelly and whom she had tried to transform into everything she wanted, she at last gets all her desires and dreams realized through passionately kissing the lifelike robot.

Rachilde's cautionary tale adds an important dimension to that of Pandora: it is not only the deception and lies of the machines that we should be worried about. Rather, relationships with AI beings or robot "people" dehumanize us too, allowing us to retreat to the most selfish, subhuman version of ourselves. On the other hand, it is the pricey aspects of real relationships—the sacrifices they require of us physically, emotionally, spiritually—that make us grow and flourish as image-bearers living in community with other image-bearers, both in times of joy but also, no less important, sorrow.

When we pursue life with other people—spouses, children, friends, neighbors—we open ourselves to the possibility that one day, death, whether ours or that of someone else, will bring an end to our relationship. After all, all human beings are defined by these bookends: we are born and we die. Dreams of substituting artificial beings for real children or spouses, however, challenges and defies these natural boundaries. In an image that should disgust us, Rachilde's Raoule simply pillages body parts of her deceased husband to create a better one, one who will never die—and will never disappoint her again. It is a twisted modern vision of "happily ever after."

In the case of Pandora, we never get more details about Pandora's marriage to Epimetheus. We do not know how the years treated them—mythology isn't really interested in documenting such tales. But we know the fruit of this marriage. Pandora's children are the curses and plagues she had unleashed upon the world when she opened the jar that the gods had sent with her for this very destructive purpose. There is a double-entendre that the original audiences would not have missed: the container is akin to a womb, and opening it was parallel to delivering a baby.

Pandora turned out to be fruitful indeed, and these children of hers are a good image to consider as a final warning. Every relationship bears fruit, but not every fruit is good. Indeed, not every fruit is human. The fruit of relationships with automata and the larger world of AI that they represent can only be something monstrous, cursed, deformed—something that seems to be an imitation of the thing, a form of it, just dangerously enough reminiscent of the real thing as to deceive some to receive it as true.

The pagan gods lied to Epimetheus—and the rest of humanity in offering this artificial being as a fitting companion. But God did not lie when He knew that only a fully fleshed human, another image-bearer, could be a fitting companion for Adam, and could in relationship teach him to be fully and truly human.

The Mad Farmer on Claude, AI, and the Church

By Hayden Nesbit

While reading a recent post on AI that gained sudden white-hot virality, one line from Wendell Berry kept coming to mind. In his poem titled *Anger Against Beasts*, Berry reflects on man's will being schooled in having its way to "transcend the impossible in simple fury." This is *the* image underlying our technological discourse, and captures the pathos of the recent post––the furious acceleration of AI advancements is now on the cusp of transcending the impossible. In other words, *something big is happening*.

This bullish (alarmist?) take on AI––whether accurate or not––announced that it will soon be an autonomous worker, incapable of error. More than that, it is making decisions that reflect real *discernment*––with not only accuracy, but taste. Such an effective substitute for cognitive labor, it is feared, will obliterate nearly all thought work into obsolescence. While transcending the impossible has long been AI's ambition, the earliest critique was simple: it just didn't compute—it didn't get things quite right. That is rapidly changing. The once clear line between AI-generated and human work is increasingly blurring.

But wherever you sit on the spectrum of concern, keeping our heads above the rising technological waters will not hinge on embracing or rejecting AI, but on embracing or rejecting our own limits. Perhaps the truest mark of humanness will no longer be perceptive finesse, but the recognition of our creatureliness—our uniquely human capacity to defy predictability.

Such creatureliness––contrary to perfection––is depicted in Berry's poem *Manifesto: Mad Farmer Liberation Front*. After opening with a sterile picture of quick profits, annual raises, and everything ready-made, Berry encourages us with this:

So, friends, every day do something

that won't compute.

This is the essence of the Mad Farmer's approach to the "everything ready-made" world. For us, as technology in general and AI specifically becomes more and more computationally sophisticated and removes unpredictability from our experience, one way to resist being swallowed by the machine is to embrace our freedom to *not* compute.

Berry gives a long list of things that do not compute––things that will seem progressively "mad" in the age of AI:

Love the Lord.

Love the world. Work for nothing.

Take all that you have and be poor.

Love someone who does not deserve it.

Ask the questions that have no answers.

Invest in the millennium. Plant sequoias.

Say that your main crop is the forest

that you did not plant,

that you will not live to harvest.

Say that the leaves are harvested

when they have rotted into the mold.

Call that profit. Prophesy such returns

Expect the end of the world. Laugh.

Laughter is immeasurable. Be joyful

though you have considered all the facts.

There is a quiet competence––an ambidexterity––to Berry’s Mad Farmer. He is comfortable moving in different contexts; he looks both near and far. He moves with a *sanctified shiftiness*––toiling, sweating, using the tools at hand, but then claiming a harvest that he never touched nor will ever see. He moves seamlessly between the present, yet somehow a thousand years ahead, embodying an agility under providence.

It is this beholdenness to providence rather than technique or outcomes that allows him to adopt an almost playful unseriousness that confounds the powers-that-be:

Go with your love to the fields.

Lie down in the shade. Rest your head

in her lap. Swear allegiance

to what is nighest your thoughts.

As soon as the generals and the politicos

can predict the motions of your mind,

lose it. Leave it as a sign

to mark the false trail, the way

you didn’t go. Be like the fox

who makes more tracks than necessary,

some in the wrong direction.

As soon as corporations have collected every byte and algorithms have mapped the ‘motions of your mind’… lose it\! Once they believe you’re squarely ensnared in their fetters of addiction, lay down false trails. Break the predictable patterns of digital life with a non-computational walk through the woods. Lie with your love in the shade. Wisely eliminate hurry. Like a fox unable to be lured by progress, Berry is giving us a parable of a life lived in the shadowlands between assimilation and fortification.

He wants us to see that both options are equally predictable. And in the face of unprecedented change, madness is gauged by predictability. True madness, then, is to avoid either predictable position. True madness looks neither like raging "against the machine" *or* putting all your hopes in an uploaded, cloud-based consciousness. Madness, instead, looks like a *holy bifurcation*. This is the kind of divided reality the apostle Peter wrote about––that Christians are simultaneously located in this age in place and time *and* in Christ. To be in Christ is to inhabit a mysterious *betweenness*, which––in a world of linear optimization––is the only true madness left.

Such a dual identity accounts for the Mad Farmer's extra tendons, enabling him to move right and left, up and down comfortably. He has "considered all the facts", not walled off in Benedictine oblivion. But he also fully expects the end of the world. Considering the compelling facts of the newest *gnosis* doesn't lead him to ecstatic anticipation of a soon-to-be-ushered-in utopia. And yet, because he is neither impressed by the facts nor bound by their proposed future, he *laughs*. He doesn't fit neatly into "user" or "non-user"; he leaves tracks in the wrong directions, each dead-end reinforcing his unpredictability.

We see this faithful flexibility throughout the Old Testament narrative as well. Glen Scrivener has called something like this the *Abraham Option*. In Genesis 14, Abraham refrains from engaging in an ongoing battle until one of his own, Lot, is taken. He then deploys his more than capable household army. He had raised them as competent soldiers, but soldiers who knew *when to fight*.

Scrivener notes that there were two things no one could ever say of Abraham:

1. That Abraham *never* engaged.
2. That Abraham *always* engaged.

Abraham was surely a conundrum to outside observers with his competence in both abstention and action. He was free to use the skills and tools of warfare, and free to reserve them for purposes he saw fit. Such is the freedom those in Christ wield with technology—freedom to be skilled users who reserve the right to opt out of a wholesale vision of techno-salvation.

Abraham's contrary refusal to be bound by dichotomies stemmed from allegiance to a higher, God-centered salvation—echoing Berry's call at the end of the Mad Farmer manifesto:

Practice resurrection.

While this is the crescendo of the Mad Farmer's campaign of contrariness, it is also precisely the underlying promise of AI and technology in general––a sort of manufactured resurrection, to bring your dreams to life, to have functional omniscience, to free up your time to live out your passion, to realize a life beyond comprehension.

But the power behind this digital resurrection is innate. AI, it is claimed, is now building *itself*. In this way, it is a sort of digital archetype for self-creation and salvation. The Christian resurrection, in contrast, is a uncontrollable resurrection that is *received*. There is nothing more unpredictable than the Spirit of him who raised Jesus from the dead taking up residence in a human soul (Rom. 8:11).

Such unpredictability––being united to another who dwells in you––is what AI seeks to smooth over: to refine its functions and realize perfection not via the life of another, but by *itself*. But true perfection (i.e. *glory*) cannot be reduced to smoothness. In fact, Byung-Chul Han argues that beauty requires death. Smoothness is the absence of the negativity of death and is, therefore, the *antithesis* of beauty. Atrophied beauty, Han says, smooths out into something undead. This undead smoothness may have an appearance of life, but it does not embody resurrection.

True glory has a grain to it. The gospel of John brings together Jesus' glory *with* his suffering. Indeed, the glory of his resurrection is inseparable from the suffering of his cross. The unpredictability of glory *through* suffering will never compute with AI. However, it is precisely in such mystery that God reveals himself to us.

Walter Rauschenbusch prayed against eyes blind to mystery “when even the thornbush by the wayside is aflame with the glory of God.” God revealed his self-existing beauty in the most absurdly unpredictable way: a thornbush aflame, yet unconsumed, with all the grit and grooves of glory. AI technologies train us to look for the smooth path, with the end of linear progress and perfection. The more we give ourselves to these technologies, we cloud our capacity to behold bushes aflame because “that’s not how bushes are *supposed* to behave”. We must train our eyes for unpredictability—such vision is sharpened by lives of contrariness that move beyond the predictable bounds of AI.

It may be the reality that soon AI far surpasses any human ability of accuracy and predictability. AI may actually replace countless experts and offer more sound advice than the most seasoned and skilled professionals.

However, if our aim is to be human, there may be no better starting posture than this from *The Contrariness of the Mad Farmer*:

I have planted by the stars in defiance of the experts,

and tilled somewhat by incantation and by singing,

and reaped, as I knew, by luck and Heaven’s favor,

in spite of the best advice.

The church is *the* community of contrariness in the world. Not because it is anti-technology, but because it is pro-unpredictable-mystery. Its sights are calibrated to resurrection, not optimization. Holy worship, sacred Scriptures, mysterious sacraments, confession and repentance––these textured elements do not compute with the smooth predictability of AI.

This contrary community cannot be pinned down by technology, but must be a congregation of image-bearing Mad Farmers who plant and till and reap in seemingly wrong directions by unpredictably glorious and mysterious means––Heaven's favor––in spite of AI's best advice.

"AI Safety" Without Virtue Will Never Be Safe

By Robin Phillips

The recent return of Sam Altman to the helm of OpenAI has brought new urgency to the nationwide debate about AI safety.

Altman is acutely aware of the dangers of artificial general intelligence (AGI), has warned that AGI could harm the world, and has been unable to alleviate fears that his company might inadvertently be building something dangerous. Yet rather than joining Eliezer Yudkowsky and the doomers, Altman has positioned himself with the accelerationists. Their creed: race to conquer AI before it conquers us.

It is widely speculated, not without some evidence, that this accelerationist approach may have been behind Altman's recent (though short-lived) ouster. Specific concern has been raised about a mysterious "Q\" *or "Q-star" system the company has been developing. By integrating existing language processing power with new capabilities in mathematical reasoning, Q* might be investing AI with capabilities for which we are unprepared.

The solution for all this, we are told, is AI safety, safety, safety. In practice, this means regulations, rules, guidelines, standards, monitoring, policies, watchdogs, etc. And here the Biden administration seems to be ahead of the game: on October 30th, the President released an Executive Order promising "new standards for AI safety and security."

Philosophers are getting into the game by adding ethics into the discussion, with the result that most philosophy courses on ethics now include discussion of AI and robotics.

Yet in all the discussion about AI ethics and safety, you'll be hard pressed to find much discussion about AI and virtue. Virtue, though enjoying something of a resurgence among contemporary moral philosophers, has for the most part receded from popular vocabulary. For many, the concept of virtue may bring to mind the quaint world of Jane Austen when characters were concerned with preserving or not losing their virtue.

What Are Virtues?

Historically, virtues have been central to the good life, and have been understood to be character traits constitutive of human flourishing. Virtues include qualities like fortitude, integrity, temperance, attentiveness, prudence, etc. Such traits are not developed overnight but require steady habituation over time.

These virtues involve more than merely knowing the answer to the question "what should I do?" but include a complex amalgam of intuition, character, and a type of know-how that emerges over a lifetime of formation, usually in community with others who are wise and virtuous.

Virtues vs. Ethics

Ethics, by contrast, is less about persons, and more about actions, duties, and laws. It is the branch of philosophy dealing, among other things, with questions like, "Under what conditions would such-and-such be morally permissible?" Whereas virtue is, by definition, personal and human-specific, ethics can be discussed in the abstract, as when we formalize the ethics in different professional fields (i.e., business ethics, medical ethics, legal ethics, etc.)

It isn't hard to see why those involved in AI are focused on ethics while tending to neglect the virtues. On one level, this is purely practical: it's comparatively easy to do ethics well, but very hard to excel at virtue. Ethics is something you can learn or teach quickly; virtue requires a lifetime of good habits and character formation. Ethics is something that can be programmed into a machine; virtue is something only humans can have.

How We Became a Post-Virtue Society

There are also historical-cultural reasons for the eclipse of virtue among AI researchers and practitioners. In his seminal work After Virtue, Alasdair MacIntyre showed that it was only after the Enlightenment that ethics and morality came to be perceived as something detached from teleology, and therefore separable from the questions of human flourishing that had been central to all previous accounts of the virtues. This led to a number of difficulties, including the classic is-ought problem made famous by David Hume: how do you infer normative claims from purely empirical observations?

In MacIntyre's reading, what emerged from these shifts was "morality" coming to exist as its own domain, suspended in a nether realm that is neither theological, legal, aesthetic, nor phenomenological, occupying a cultural space entirely its own but with a rather hazy ontological status. By thus detaching morality from the rest of life, the philosophers of the 18th century created space for interminable questions about whether morality is even real (or, as the problem is often framed, whether morality is "objective").

Beyond conundrums about the objectivity of morality, the philosophical sea-change of the Enlightenment bequeathed a legacy in which there is no rational context in which to situate the virtues. Consequently, we live in what MacIntyre described as a post-virtue condition, in which the status previously afforded to the virtues has been replaced by what he calls emotivism. In MacIntyre's analysis, virtues remain a driving force in our public discourse, yet the virtues to which we might appeal (whether tolerance, justice, cooperation, equity, etc.) exist in a merely fragmentary form, the remnants of an older tradition that once provided a coherence now lacking.

What Has Virtue to Do With AI?

Nowhere is this post-virtue condition more evident than in the field of AI. Those working in the burgeoning domain of "AI safety" are proceeding on the assumption that virtues can be neglected in favor of a reductionism in which AI ethics is simply one more programming problem.

We see this reductionism in OpenAI, the company responsible for ChatGPT. For example, in his AI safety lecture, Scott Aaronson from OpenAI suggested that programming goodness into AI is being approached as a problem of engineering. Aaronson shared how he was under pressure from Ilya Sutskever, cofounder and chief scientist at OpenAI, to find "the mathematical definition of goodness" and "the complexity-theoretic formalization of an AI loving humanity." Here is perhaps the final culmination of Hume's is-ought problem: what is real is what can ultimately be reduced to algorithmic mathematics.

How might the discussion of AI safety be reframed if we positioned the virtues as central to the debate? I suggest this would open up new sets of questions currently neglected in the industry, including perhaps:

- What habits would engineers and managers need to cultivate to steer tech companies (and by extension, the industry) in a humane direction?
- What might an AI company look like that valued engineers and managers with virtues like moderation, prudence, temperance, courage, etc.
- What sort of character traits should be prized among those working to align AI with human values?
- What types of habits of mind and body might be conducive to the cultivation of such traits over time?

For many, these types of questions may seem a category mistake. When companies like Microsoft and Google are hiring engineers, the last thing they are likely to ask is whether the candidate makes his or her bed every morning or practices prudence in relationships. In fact, what stockholders seek in corporate executives is not virtue but skills. As long as executives do not break the law or violate professional ethics, their level of personal virtue is seen as irrelevant. Similarly, the stereotype of the good programmer is often precisely someone who lacks human virtues, perhaps a person on the spectrum whose brain works in a single-minded, Spock-like fashion—in short, someone like Sam Altman.

As this neglect of virtue becomes systemic to an entire industry, the result is what we see today: largely amoral tech companies concerned primarily with making money, with ethics departments obsessed with laws, protocols, guardrails, codes, policies, etc., but with comparatively less attention on the virtues.

President Biden's Executive Order on "Safe, Secure, and Trustworthy Artificial Intelligence" is a case in point. The sprawling 19,682 word document never mentions virtue, but treats AI safety and ethics (often grouped under the larger umbrella of "trustworthy AI") as problems that can be adequately addressed through rules and guidelines alone. Of course, law cannot mandate personal virtue, but that is precisely the point: if we trust the law alone and neglect the larger conversation about virtue, we will be opening ourselves up to a false security.

Beyond Law: A Plea for Human Virtues

The current law-based approach echoes the ethics of Immanuel Kant (1724–1804), whose answer to Hume's is-ought problem was to place ethics on the same footing as scientific laws. For the Prussian philosopher, ethics became a discipline dealing with impersonal laws that are universally descriptive—what he famously called the "categorical imperative." By thus detaching morality from the personal, the human, and all consideration of circumstance, Kant imagined he had rescued ethics from receding into subjectivity. On such a scheme, ethics does not emerge out of the messy and imprecise realm of virtuous habits, let alone character traits developed over time (Kant actually downplayed the role of character); rather, ethics involve knowing the right ethical laws and then going out into the world to apply those laws irrespective of circumstance. As such, ethics becomes not a matter of situation-specific prudence, but something like the law of gravity, an exact science dealing with impersonal laws.

To any of the ancient wise men—Confucius, Plato, Aristotle, Jesus, the saints of the Christian tradition—the idea that ethics could be impersonal would have been nonsense. For them, ethics is about nothing if not character. Even the first five books of the Hebrew scriptures, which go a long way to itemizing laws to cover almost every conceivable circumstance, ultimately contextualize ethics in terms of character traits and wisdom that emerge through a well-disciplined life over time.

It is significant that contemporary AI safety works, not in the wake of ancient philosophers or the Bible, but in the shadow of Kant. For example, Isaac Asimov's seminal (though largely hypothetical) work on robotics in the 1950s proceeded on the assumption that robotic safety could be achieved by programming machines to follow the right procedures. According to this framework, if something goes wrong in an automated system, whether in real life or in a simulation, then we simply need more comprehensive laws. Yet as the AI community is finding out from the notorious alignment problem, such an approach to AI safety may not be feasible within current, or even future, engineering frameworks.

The turn to human-specific virtues, though it strikes many as inefficient, may actually prove the only option in the era we are quickly entering. Yet our society seems largely unready for that discussion. Modern tech companies like Microsoft, Google, Facebook, and OpenAI, seem intent on continuing the legacy of the 19th century scientific management movement through the fetish with "efficiency,"—a word which, in practice, functions as code-speak for achieving all but total automation many degrees removed from human agency. Through AI-powered automation, it is hoped we can eliminate the variabilities and unpredictability of "human factors," and thus achieve the technological analogue to Kant's categorical imperative.

Yet we would do well to pause and consider the debt we owe to human factors, especially a type of virtuous intuition that cannot be quantified. Consider the

1983 near-disaster when the Soviets nearly launched the world into nuclear holocaust. Around midnight on September 26, a time of particular strain in the Cold War, unusual weather patterns caused the Soviet early-warning radar system near Moscow to incorrectly register five incoming American missiles. The computer responded by issuing a retaliatory order "Launch," which would have triggered a full-scale nuclear war against the United States and her allies. Even after procedures were followed to rule out computer error (for example, by resetting the systems), the order to retaliate continued.

Stanislav Petrov, the lieutenant colonel on duty in the Soviet Air Defense Forces Dome that night, was cool-headed and thoughtful. Petrov drew on his past training, which had convinced him that any U.S. first strike would be massive, likely more than only five missiles. While the rest of the base was in panic, Petrov put his career on the line, deviated from orders, and declined to follow the instruction to launch an attack. That night Petrov exercised the virtues of patience, prudence, attentiveness, thoughtfulness, and courage, while avoiding vices like reactivity, impulsivity, rigidity, and fear. As a result of his virtue, a nuclear disaster was averted.

If this incident had happened not in 1983 but in 2023, the outcome would likely have been very different. Today widespread bias for automation has resulted in a shrunken space left to human virtue, including situation-specific prudence. One of the reasons that systems regulating everything from commerce to trading to military defense are increasingly handed over to AI is precisely to eliminate the messiness of human judgment, as if "human factors" are a bug and not a feature. Yet without human agency there can be no virtue, including the type of virtue Petrov exercised when he proved the triumph of the human over the machine.

Discussion of possible AI doomsday scenarios are often compared to the nuclear threat. Yet if we can learn anything from our near brush with nuclear holocaust, it is that rules, guidelines, policies or procedures will never save us: we need old-fashion human virtue.

New Calvinism in an Age of Enchantment and AI

By Cameron Shaffer

The New Calvinist movement needs renewal, not through innovation but the rich sacramental theology of Old Calvinism. As enchantment returns and AI emerges, these deeper resources of the confessional Reformed tradition are an indispensable missional and gospel-centered gift.

New Calvinism is a movement fostering "gospel-centered" partnerships in a Reformed-Protestant key across evangelical churches. To be "gospel-centered" means that the movement prizes the doctrinal clarity of the gospel from a Reformed perspective and then prioritizes the clear communication of that gospel (in preaching, apologetics and evangelism, discipleship practices) within and outside the church. Despite the name, New Calvinism is largely shaped by Jonathan Edwards and his theology of affections.

As a Presbyterian minister committed to a Reformed ecclesiology, I see this informal coalition as a net positive. The kingdom of God is more than Presbyterian and Reformed (P\&R) denominations, and crossing denominational lines for the sake of a clear and centered gospel ministry is a good thing.

Yet the movement needs refreshing. 2025 marks the 20th anniversary of The Gospel Coalition (TGC; the flagship New Calvinism organization) and the "The Young, Restless, and Reformed" movement is no longer characterized by resurgent energy. Jake Meador observed that the movement attempted to capture the evangelical center, only to encounter the fracturing of evangelicalism. Similarly, Aaron Renn recently reflected on the movement's maturation and accompanying diminished influence outside its own adherents.

An exemplar of needed refreshment is TGC's "Theological Vision for Ministry," a now 20-year-old statement on what gospel-centered ministry looks like in the present. But 2005 was a long time ago and the world has changed. One of the strengths of New Calvinism was that it kept the gospel foregrounded while thinking clearly about doing ministry amid the complexities of the modern world. But that becomes a weakness when visions are for ministry in a world that is two decades out of date.

In a significant development towards this end, a large summit of New Calvinist pastors gathered at Spanish River Church in February to initiate a renewed Evangelical-Reformed (i.e. gospel-centered) cooperative mission movement. Spanish River Church, a congregation of the Presbyterian Church in America, was a crucial founding church in the Acts 29 network and a catalyst for the start of New Calvinism. The summit cast a concrete vision of how to address the changing cultural landscape with a reenergized gospel-centered ministry.

This is where Old Calvinism comes in.

"Old Calvinism," meaning the historic P\&R churches, is not just characterized by Reformed soteriology, but also sacramentology, ecclesiology, and liturgy. Old Calvinist churches hold that the sacraments are means of grace: effective means of salvation. That by the blessing of Christ, the power of the Holy Spirit, as they are received in faith, the sacraments confer the grace and reality that they represent, which is Jesus Christ himself. In baptism the recipient is ingrafted into Christ, has their sins washed away by his blood, and is regenerated by the Holy Spirit. At the Lord's Supper the worthy partaker feeds on the true body and blood of Christ and thereby receives him and his benefits.

This is the confessional language in the Westminster Standards (Presbyterian) and the Three Forms of Unity (Reformed). The sacraments, as signs and seals of the new covenant, are central to the gospel: by them the grace of the gospel, which is Christ himself, is displayed, applied and effected, and received by faith. The sacraments center the gospel in the life of the church and in the world. The new covenant is in the body and blood of Christ, and the Great Commission foregrounds baptism; these things are central to the historic Reformed understanding of the gospel. P\&R types should be unabashedly upfront about this, not just in our doctrine, but in our gospel proclamations.

While P\&R pastors and churches still held to their convictions, the New Calvinist coalition intentionally backgrounded the sacramental and ecclesial dimension of the Reformed tradition, treating them as secondary or tertiary issues to hold the movement together. Even though the movement was Reformed, the cross-denominational coalition privileged Baptists over historic P\&R churches, due to a combination of sheer numbers and Presbyterians being willing to downplay some doctrinal distinctives for the sake of missional partnership. For instance, TGC's "Foundation Documents" have minimal references to the sacraments, and what little they have are either broad statements about what they represent, or, in the single passing reference in TGC's "Theological Vision for Ministry," are included in a list of formative practices. Likewise, Tim Keller's textbook on gospel-centered ministry, *Center Church*, relegates the sacraments to formative community practices and opportunities to proclaim the gospel.

Yet in confessional Reformed theology the sacraments are not incidentals to gospel ministry, but central to gospel ministry. The sacraments are more than "visible words" which proclaim the gospel. With the sacraments, God not only speaks, *but God acts*. The sacraments are not just reminders or challenges related to the content of the gospel message, but a divinely established avenue by which the gospel is applied to the heart of the partaker so as to conform them by the Holy Spirit into the image of Christ.

Presbyterians in the gospel-centered coalition should press for a re-centering and prioritization of Reformed sacramentology as part of a recalibrated New Calvinist-evangelical collaboration — not as something annexed to gospel-centered ministry, but as part of its core identity. Yet in Spanish River's renewed call for gospel-centered ministry the sacraments receive a single passing reference as things that should be administered in a healthy, biblical church. This is not renewal, but rehashing.

Besides doctrinal conviction and wishful thinking on my part, why should this recalibration happen? The gospel-centered movement has always required a missional impetus to justify emphases, and relitigating old doctrinal debates doesn't fit that bill. And there are at least two big reasons: re-enchantment and AI.

Re-enchantment, a thesis that has been the subject of much discourse, holds that declining traditional religion in the West has been replaced not with a materialistic secularism, but an increased belief in "weird" spiritual and supernatural realities whose viability had been suppressed by modernity. As Christian Smith demonstrates in *Why Religion Went Obsolete*, re-enchantment is a significant factor in the decline of traditional religion, which appears to inadequately address the phenomenon. This began with Gen-X and is increasingly true for each subsequent generation.

Reformed theology understands the sacraments as being fundamentally spiritual. The sacraments are not just divinely inspired object lessons, but the physical being united to the spiritual and us receiving it. And the sacraments are not just one supernatural talisman among many, but signs and seals of the true God who comes to us in the person of Jesus and the power of the Holy Spirit. This is truly spiritual, and no other supernatural experience is actually better.

The missional opportunity to proclaim the gospel of the union of heaven and earth, begun in the incarnation of Jesus as he fills creation with his Spirit, testified and conferred in the sacraments of his body and blood until he comes again in glory, is immense.

In a striking moment at the most recent TGC conference, John Piper read a prayer composed by ChatGPT in the voice and theology of D.A. Carson. The composition was quite good. Piper's critique was that there was no feeling or praise in it. That is a strong and consistent Edwardsian, New Calvinist reaction. Whatever "thinking" the AI may have done, it has no affection for God. But as an explanation of what makes people unique, distinct from any sort of artificial intelligence, it doesn't altogether work.

Perhaps AI will never truly think, feel, or will; perhaps that's impossible. But you don't have to believe in the predictions of AI 2027 to recognize that the things we once thought were the exclusive domain and character of humans may be possessed —or at least mimicked in such a perfect facsimile as to be indistinguishable from human thinking, feeling, and willing —by AI.

The question of "what makes humans different and meaningfully distinct from AI?" is going to become pressing and existential. And the answer is not thinking, feeling, or willing, whether or not AI actually surpasses humans. The answer is that only people created in the image and likeness of God can have faith in him, be united to Christ, and have communion with God. Only humans can participate in the divine.

Perhaps one day AI-powered robots will possess consciousness and be capable of eating. But no matter how advanced they get, even if they eat the bread and drink the wine, no AI machine will ever be able to eat and drink the body and blood of Christ. This is the 21st century update of Aquinas' teaching that mice and dogs may be able to eat the physical substance of the sacrament, but never its true substance, who is Christ. No bot will ever be indwelt and baptized by the Holy Spirit. The missional impetus for centering the sacramental dimension of the gospel is that they are the visible and tangible deliverance of God to man, something that no AI could ever give or receive.

This is central to the mission and ministry of the Church. The gospel is that we are redeemed by Christ for union with God in the life of his Spirit. This redemptive union and life is signed, sealed, exhibited, and conferred in the sacraments. The nature of the sacraments should be more important to New Calvinism than the questions of who gets baptized with how much water and how often the Lord's Supper is administered. So what could an Old Calvinist, gospel-centered-sacramentalism look like in a refreshed New Calvinist movement?

First, the obvious: P\&R members of the movement should press for Reformed sacramentology to be prioritized in doctrinal statements and ministry philosophies. When the movement discusses being gospel-centered (preaching, apologetics and evangelism, discipleship practices) Reformed, confessional, sacramental theology should be included. On a practical level, centering the sacraments with a Reformed understanding allows an Evangelical-Reformed alliance to expand its denominational boundaries for the sake of the kingdom. In the same way that New Calvinism itself was initially an expanded, cross-denominational movement, Reformed sacramentology allows the movement to expand again.

Second, the church should clearly talk to its people about what the sacraments are and what God does through them. We should stress with Martin Luther that we are the baptized ones: We are in Christ because we are baptized into him. This is critical to discipleship. We are washed, we are filled, we are Christ's. This is the Westminsterian idea of "improving" our baptism —we grow into who we truly are by relying on the Spirit who has been given to us as we rest upon what Jesus has done for us. When we come to the Table we have true communion with God because he grants us "so to eat the flesh of thy dear Son Jesus Christ, and to drink his blood, that our sinful bodies may be made clean by his body, and our souls washed through his most precious blood, and that we may evermore dwell in him, and he in us."

Finally, to the world, our gospel-centered evangelism should have a sacramental dimension. Would you be cleansed? Would you be filled with the Spirit? Would you pass from death to life? Be baptized into Jesus in the name of the Trinity. We should unashamedly maintain the Petrine order in our gospel proclamation: Repent and be baptized for the forgiveness of sins and you will receive the gift of the Holy Spirit. Would you partake of the divine? Would you receive the transcendent and most real? Be redeemed and united to God? Receive the body and blood of Christ. Doing this would refresh and strengthen the New Calvinist movement for ministry in the years to come.

If You Ask A.I. for Marriage Advice, It'll Probably Tell You to Get Divorced

By Samuel James

How to watch a massive cultural shift as it happens: A story in two parts.

Part I

A user on the online forum Reddit used data analysis tools to study the kinds of comments people leave for people asking relationship advice. The results are both alarming and unsurprising.

Reddit's reputation in this regard is not a secret. There are entire personal essays out there about the swiftness with which Reddit users will urge questioners to blow their relationships up. In 2022, the lifestyle blog Refinery 21 published an article titled "The Interent Said 'Dump Him…' So I Did." Most of the stories covered in that article cover things like cheating or abuse—the kind of behavior most people would likely urge someone to walk away from.

But Reddit's relationship advice sections often fail to distinguish between things like that and other conflicts. For example, when one user recently complained about her boyfriend's self-pity after her medical procedure, the comments blew up. A sample of the top rated replies:

"He is so beyond pathetic here"

"Has he always been stupid or is this new behavior?"

"This guy is a self centered dirtbag. What an AH to not even care about you and your recovery. It was only what could you do for him. That's so gross. Dump this \[expletive\]."

"He will never be there for you. A huge red flag."

Reddit's relationship advice forum prioritizes simplistic answers that emotionally validate the user. There's often no sense on Reddit that relationships are worth fighting for, possibly even worth enduring some pain for. If the significant other is not meeting your expectations, your best bet is to end things right there.

That has always been true of Reddit. But look at the data chart again. Since 2010 in general, and since 2020 in particular, the percentage of "end the relationship" comments has skyrocketed. And the percentage of comments encouraging communication or compromise has fallen significantly. If you ask Reddit's relationship advice forum what you should do today, you are, statistically anyway, about 3x more likely to be told to end the relationship as you will be told to communicate.

In my book *Digital Liturgies*, I explain how the logic of Internet life makes these kinds of relational blow-ups seem more plausible. Digital habitats are immediately responsive to our wishes. We can delete, mute, block, and un-friend at will. As more and more of our lives are mediated through digital technology, our emotions become trained to expect this kind of responsiveness everywhere in the world. We feel like something is wrong if another person is allowed to make us uncomfortable and we cannot "exit" the situation. In that way, the Internet lies to us about what real life is like, and develop an inability to live with real people in real relationships—which are often difficult.

I'm convinced that part of the emerging polarization between men and women has to do with the increasingly niche information streams that men and women are immersed in. Men see the excesses and abuses of feminism daily. Women see the excesses and abuses of masculinity daily. Modern men and women are caught in de-sympathetic rhythms of polarized gender discourse. Online male culture and online female culture distrust each other, in part because online life highlights relational tension as a content currency, and also in part because marriage and child-raising, which build empathy between the sexes, are declining. Men and women know less about each other right now, and what they do know is often negative.

Reddit's user base reflects this. While it's true that the kind of cases that would inspire people to write online about them probably skew toward more extreme situations, it's also true that Reddit's relationship advice forum is very large, and very active, and covers an immense spread of scenarios. The online dynamics mean this data is not surprising. What is alarming is the acceleration of relationship-ending advice relative to other, more careful pieces of advice. This acceleration matters because…

Part II

…it turns out that AI companies are leaning on Reddit's content to train their bots. Axios reports that major A.I. bots frequently give answers to user questions that pull from Reddit content. In fact, according to Axios, Reddit is actually the \#2 source of content for LLMs (large language model, like ChatGPT or ClaudeAI), behind only YouTube. Here's a sobering quote included in the Axios report:

At last month's Tatari conference, Reddit CEO Steve Huffman summed it up with two slides: "Today's Reddit conversations are tomorrow's search results" and "No matter how good AI gets, people will always want to hear from other people."

Practically speaking, this means that users who ask AI bots for counseling or therapy—which is right now a lot of people, and is going to be a lot more people in the future—are going to get a lot of answers pulled from Reddit. In other words, these LLMs are going to spitting out answers to questions like, "Should I get divorced," by repeating how users on Reddit answer those kinds of question.

And we know how users on Reddit tend to answer those questions\!

One of the things worth mentioning about AI counseling is that these bots skew heavily toward validating the user and giving him or her the answer AI thinks they want to hear. There are inherent biases in AI programming that steer it toward sycophancy and a customer-is-always-right ethos.

This is a powerful influence on AI advice or therapy. If a user asks ChatGPT for relationship advice, the AI is already primed to deliver what will satisfy the user most. If ChatGPT is trained on Reddit posts, it will believe (via statistics) that the most likely thing to satisfy the user is to counsel ending the relationship. This won't be true in every situation, but remember, AIs do not think. They simply repeat input. Any computer that learns about users' relationships from Reddit posts is going to conclude that ending a relationship is the most preferable option, because that is what Reddit commenters say.

I'm not trying to be overly doom and gloom, nor am I denying that AI bots might counsel something better than divorce or breakup. They can, do, and will. But what seems inarguable at this point is that looking toward the Internet for relationship advice is a good way to get one type of advice over and above anything else. This represents a very serious chance at a massive cultural shift. Technology shapes how we think, and it's not every day that we get to see behind the curtain as this power is cooked up. But we're seeing it now.

And here's where I encourage any pastors, counselors, or youth leaders who may be reading this:

AI is going to be a major player when it comes to therapy and counseling. Similar to the way Google has transformed how we look for medical advice, AI is going to play the role of sage for a lot of people. There's a threat and opportunity here. The threat is that Christians will be sucked into destructive decision making as they consult what they naively believe are objective "data" machines. The opportunity is that AI, like all digital tech, is very unsatisfying. Just as porn doesn't satisfy intimacy and inboxes don't solve loneliness, AI will not leave users feeling cared for.

That's something only people made in the image of God can do.

Artificial Intelligence is a Technology

By Matthew Schultz

A pair of dice rolls across the table. A man glances at a chart and then pulls a measure of music from a pile. He rolls the dice again and selects another measure. Then another. His friend watches and laughs as the piece grows longer and longer. The man sits down at the piano and begins to play. As the randomly arranged segments begin dancing through the air, his friend's face turns from mirth to furrowed brow. There is something delightfully scandalous about it.

This is the Musikalisches Würfelspiel, an eighteenth-century parlor game traditionally associated with Mozart. The trick was simple: a set of cleverly composed musical measures that would produce pleasant music no matter how they were arranged. The game appeared to reduce the musical craft to a random mechanism. Much of the discussion around artificial intelligence treats the use of the technology as a kind of würfelspiel.

Standard accounts worry users will become nothing more than curators who arrange pre-composed parts, re-rolling prompts until they get an output that vaguely matches the intention and stitching the results together, all without any attention to the slow, careful development of the fundamental skills essential to craft. And the question of how artificial intelligence relates to craft is important not only because the work that we do is constitutive of a functioning society, but because many forms of work are so integrated with thick community relationships and practices that disruption in these spaces can be close to identity-destroying.

While this is an intuitive picture of the technology and its applications, this evaluation is stuck a few years behind the actual practice. Earlier iterations of frontier models were crude: users would enter a prompt, press enter, and be greeted with text, image, or audio, much like rolling a die and "composing" a song. There was limited control of the process or the outputs, and there was a glaring disparity between the amount of effort from the user and the quality of the output. However, modern workflows have become remarkably more complicated over the last few years and can be most clearly seen in two examples: image generation and music production, where workflows are highly granular and involve a great deal of control over the creative process.

One popular local UI framework, ComfyUI, allows the user to deploy hundreds of custom nodes to finely tune inputs and outputs at every stage of generation. Advanced workflows using such tools begin to resemble something like a synthesizer board, with different dials and knobs tuning the fundamentals of the image: subject, composition, color, framing, and style. A user might draw a sketch and create a node setup that reworks the line art, adds specific pieces and styles of clothing, and then imagines the character on a turnaround sheet, mimicking a concept photoshoot. (It should be no surprise that these sorts of workflows are already being incorporated into the fashion industry.)

Music production has seen similar shifts. Large numbers of musical artists and producers have integrated Suno, a premier music generation model, or similar artificial intelligence programs into their workflows. Inexperienced artists who simply enter a prompt and publish the results find that their work is largely (and rightly) ignored. At the other end of the skill distribution, the most advanced artist workflows are similar in complexity and granularity to image generation and involve significant amounts of human input based on knowledge of music theory and sound production. A sample can be produced by an AI program according to specific parameters, its stems separated out, and the results dropped as separate tracks into professional studio software. This is then handled as a standalone sample, remixed, or otherwise edited and incorporated into a full song with a synthesized drum track and human-composed instrumentals or vocals.

The most telling development across these creative spaces is that most of the concerns with machine learning seem to terminate in discussions around the ethical sourcing of data sets rather than suggestions that the technology is fundamentally outsourcing creative capacities. Disruption here is already turning into creative adoption, which is what we would expect if the technology was another form of acceleration continuous with the industrial revolution rather than something completely different. The fact that these communities can distinguish between a deskilled and skilled use of the technology cuts directly against deskilling worries as totalizing or novel. Practitioners are applying and developing core fundamentals to new technological processes and discussing and exploring successful integrations in ways that begin to look suspiciously like a craft.

However, even if we grant that modern practitioners are able to exert fine-grained control of the models, this might ultimately be a shallow control given that we still don't understand how these models work at a fundamental level. It is reasonable to worry that a concern about implicit knowledge in traditional craft can naturally develop into a full-blown rejection of artificial intelligence if we continue to use a tool we could never understand. Auditing a model's internal processes to improve outcomes is a genuine concern that seemed, at least for a while, to represent an impassable frontier of knowledge.

Yet the mysterious nature of these models will not remain mysterious forever. Researchers have begun opening these black boxes and mapping their internal structures. For example, current research has discovered that the internal model of Claude counts character lines using 6D helical manifolds, a geometric approach to variables that mimics place location neurons in mammals. This suggests the internal structures of these models are elegant and orderly rather than illegible. That computer scientists are beginning to reverse engineer the technology in a similar manner to how we have reverse engineered the brain should not be surprising. As Christians we live and work under a cultural mandate that not only presupposes creation is knowable but charges us with the task of knowing it.

If AI is not *sui generis* with respect to its positive practices, neither is it with respect to its dangers. Some of its uses will deserve to be categorized under an umbrella of technologies designed to hook directly into your dopamine systems to maximize profit. Furthermore, even setting aside workflows where AI is doing material work—protein synthesis, diagnostic imaging, drone-directed agriculture—most use cases tend to neatly integrate workflows that exist entirely in the digital space. AI is unlikely to restrain, let alone reverse, the most common digital modes of production.

Yet the greatest danger is both more pervasive and less obvious: AI is much more likely to be deployed as a multiplicative layer that allows ever more efficient micro-targeting of digital services and physical products by industries that already profit from compulsive behavior. The advent of hyper-personalized, real-time engagement strategies will require legislative safeguards, especially if AI leads to video advertisements generated in real time for an exhaustively mapped individual profile.

The question of how this intersects with a Christian anthropology is necessarily complicated. Much of what has been written exists at a theological level with a defensive posture, targeting the fringe ideology of Silicon Valley tech elites in an apologetics framing or treating the use of AI as an unwitting portal to the demonic. Yet if the continuity thesis holds, then we should ground our response in a few important fundamental remembrances.

Technology has existed since the Garden and is an integral component of our cultural mandate. We should also remember that one of the core distinctions between the Creator and his creatures is that we never create matter but merely (\!) rearrange it. This becomes clear whether we consider an ancient farmer in Mesopotamia irrigating a plot of soil, a medieval peasant in Northumbria weaving a basket from flax, or a young musician in London taking the raw outputs of machine sound, adjusting its pitch, volume, and length, and incorporating it into a DAW loop. While there are all sorts of important distinctions and qualifications between pre- and post-industrial craft, there is no metaphysical distance between the two.

If the continuity between industrial technological disruption and its digital age permutations holds, then so do the spiritual realities. As Joseph Minich demonstrated in his book *Bulwarks of Unbelief*, our experience of the post-industrial world is phenomenologically bent toward unbelief. Rather than breaking with this posture, artificial intelligence represents an acceleration of this pressure. This suggests not a novel response, but a deeper one consistent with those practices that have always defined Christian community.

I am thinking here not only of my church's weekly dinner and hymn sing but an elder who is also an artist who is working on a mural painting for the narthex, a leader practicing his craft in a local community. These sorts of practices create the ground—the posture of learned humility, embodied friendship, and grounding in a historical tradition—needed to cultivate people who can use technology responsibly. Our task is not to develop a unique theology of AI but to catechize our members into a people who can wield this technology without becoming captive to its internal logic. Like alcohol, artificial intelligence will become a test of character, a dangerous good that divides the foolish from the wise.

Wendell Berry’s Unanswered Question

By Samuel C Heard

Earlier this year, the language education company Duolingo caused a stir when its CEO, Luis von Ahn, put out a memo announcing his intentions on making Duolingo an “AI-first” organization. He wrote in April:

“AI is already changing how work gets done. It’s not a question of if or when. It’s happening now. When there’s a shift this big, the worst thing you can do is wait. In 2012, we bet on mobile …. We’re making a similar call now, and this time the platform shift is AI.”

As part of that announcement, von Ahn told his employees that they would begin to phase out contract workers, explaining that they were doing the kind of work that could be handled by AI: “Being AI-first means we will need to rethink much of how we work. Making minor tweaks to systems designed for humans won’t get us there.”

It turned out that not everyone was thrilled about these changes, and Duolingo faced a considerable amount of backlash for the announcement. This apparently came as a surprise to von Ahn — he wasn’t sure why people were targeting Duolingo, seeing as other companies around him were making the same decisions. “Every tech company is doing similar things, \but\] we were open about it,” he told the [*Financial Times*.

In this regard, von Ahn has been right. Though not nearly as outspoken as Duolingo, over the past year a number of other companies have indeed been phasing out workers in a pivot to AI:

- Earlier in July, Microsoft made the decision to eliminate 9,000 positions, after CEO Satya Nadella recently shared that 30% of its code is now being written by AI tools.
- After cutting 1,000 positions earlier this year, Salesforce CEO Marc Benioff recently said that AI technology currently accounts for 30% to 50% of the company's work.
- IBM recently cut 8,000 employees, mostly concentrated in Human Resources, to welcome in AI agents to complete tasks.

The list goes on. Some companies seem to be pushing AI as far as they possibly can go. In May, finance company Klarna's CEO Sebastian Siemiatkowski said that the company had dismissed 40% of its workforce — including the entirety of its customer service department — so that they could prioritize AI implementation. The company even touted a "replacement" of Siemiatkowski himself by using an AI-generated version of him to give a quarterly update. It's no wonder von Ahn was baffled.

Here we face the irony of our present moment: A company like Duolingo faces backlash for going "AI-first" just as the novel technology gets stitched into the fabric of our lives everywhere else. Not only are people now losing their positions to AI — even the workers who have kept their jobs are now being regularly briefed on how to implement AI and LLMs in their line of work. Browse any major website, and you're likely to find a new "AI tool" to help you carry out your tasks. AI-generated texts now populate our search results before finding links to other pages. We're having trouble discerning whether that post, that image or piece of music has been created with AI. The onslaught of an "AI-first" culture is not a question of if or when. It's happening now.

And yet, our criticism of companies like Duolingo for making this culture explicit shows a state of dissonance and unease about this whole situation. We feel conflicted; we are divided among ourselves. We don't like it when a company says it's going "AI-first." But then again, why is every institution getting involved in the technology? Why are dozens of companies firing workers to be replaced with AI? Why is every website asking us to interact with their new chatbot? If, collectively, we find the concept of an "AI-first" culture off-putting, then wouldn't our best course of action be to limit this technology's usage in everyday life?

Of course, the answer to this question is that some people do indeed want to create an "AI-first" culture — von Ahn's memo proves the point. But I wager that they stand in the minority. The rest of us remain unconvinced. Even those who regularly promote its purported benefits — the AI optimists among us — would likely shudder at the idea of an "AI-first" culture. Many, if not most of us do not want this.

But what's harder, I think, is for us to explain *why* we don't want this. And there's the rub: No matter how much we dislike it, an "AI-first" mindset can still grow in a culture that largely feels opposed to its domination. All it takes is a community that cannot properly justify why its work must be done by *people*. A culture that has no stable footing in its own worth.

Why Should People Do Our Work?

The simple truth is that it's actually pretty straightforward to resist the growth of an "AI-first" culture, if that's indeed what we want. All we'd need to do is answer the following question: "Why should people do our work?"

On the surface, even that sounds a little foolish. "Why should people do our work?" feels akin to asking why we would choose the sun for daylight. What other choice do we have?

But of course, the past few years have shown us that we do indeed have a choice between a person and a non-person to accomplish our work for us. If I'm looking for a piece of art for my living room, I can now prompt an image to a chatbot just as easily as I could commission an artist or photographer. An employer deciding whether to hire a copywriter must now factor in the reality that at any moment, an LLM can produce the necessary writing on demand. People experiencing mental health crises are no longer required to seek out counselors in person, but can now interact with AI to diagnose their struggles. It turns out that the sun is not our only option for daylight. Artificial light sources abound.

By its very definition, an “AI-first” culture is one whose members collectively and repeatedly choose the artificial “non-person” over the person. So it follows that, if this is the sort of culture we’re trying to avoid, then we need to choose the person over the non-person in the majority of our interactions with the world. Maybe there is a place for AI chatbots in our ideal culture. But this technology would need to remain subservient to the actual choosing of persons if we want to avoid going “AI-first.” In other words, if given the choice between a graphic designer and an AI model, then we need to be willing to choose the designer, at a minimum, 6 out of 10 times — or better yet, 9 out of 10 times, or even 99 out of 100 times.

Hopefully you see where this is going. To choose a person over a non-person — a graphic designer over an AI chatbot, for example — looks nice on paper. But pause to consider what this actually means. In choosing the designer, you have opted for the following realities:

- The designer might disagree with you or feel resistant to your wishes. (AI has been engineered to agree with you.)
- The designer’s work takes time, sometimes days. (AI produces images in seconds.)
- The designer is a human being, and that means feelings are involved. It takes social energy to engage this person face to face. (AI demands nothing from you.)
- Most importantly — the designer requires a paycheck. (AI requires, at most, a subscription.)

Now, our question becomes a little more clear: Why, exactly, should we rely on *people* to do our work? After all, it's hard to compete with a machine that "doesn't go on strike" and "doesn't ask for a pay raise," as one CEO recently put it. These things produce images and words instantaneously. They don't talk back, and they demand nothing other than your ability to generate a prompt. We know that our resistance to going "AI-first" means choosing the person over the non-person, but it's in our attempts to justify this choice where things start to go awry. After all, why would we want to choose the less convenient option?

Some have tried to argue for people on the basis of our skill and quality of work. At the moment, it's a popular move among many AI skeptics to claim that AI technology is not actually as intelligent as advertised, and thus cannot be trusted when compared to the quality of craftsmanship from a human being. We ought to choose the person, then, in order to get a finer resulting product.

It's easy to see why this argument is being made — when an AI company CEO proclaims its model has reached "PhD level" knowledge, and then a few days later that very model fails to recognize how many b's are in the word "blueberry" you're left scratching your head a little. Never mind the fact that AI has already made itself notorious for generating convoluted "slop" images. And of course, we all know that it can produce an email, but people are getting better at recognizing an AI-generated tone of voice.

As someone who is fairly skeptical of this technology and its usefulness at the present moment, I am sympathetic to this argument being made. But I have several concerns with this line of thought. For one, while it may be true that our AI chatbot is not a fully trustworthy guide at present, what happens if it reaches that proclaimed "PhD level" in five years? Two years? Should it evolve to become a superintelligence — the closest thing we can get to omniscience — would we still be able to justify choosing the person over the non-person? We have to grapple with this technology not only as it currently stands, but also based on what it has the potential to become.

Ultimately, as tempting as it may seem, we cannot ground our argument for people on the basis of quality of output. It may work now in some contexts, but the argument falls on its face as soon as AI creates a product of similar quality to its human counterpart. And we don't have to wait years to see this happen: Already, LLMs are crafting blogs, emails, and social media posts, all of which are usually engaging, or at a minimum coherent. Can a marketing writer create better copy? Of course — though even now, this isn't guaranteed. Does it even need to be better, though? In the eyes of marketing leads and business owners looking at the bottom line, a "good enough" text made by a cheap AI model may be a lot more tempting than hiring a writer. Especially when we set efficiency and profit as our baseline standards, no person is going to get the upper hand over the machine. We might as well throw in the towel.

But if we really do want to resist going "AI-first," then we still need to be able to explain why we would ever choose *people* to do our work, especially when we have the option to offload our jobs to machines. This proves a challenge, especially when we only rely on the quality of our products to do the talking for us.

The only other option, then, is for us to move away from these considerations of our efficiency and productivity and get to the real root of the matter. Questions of why we should rely on people to do our work — indeed, why people should even work at all — are necessarily intertwined with our purpose: who we are, why we exist, what our labor is even for. The first question we should be asking, then, should not be "Are people more efficient?" or "Can people create a better product?" Instead, we must ask ourselves what writer and farmer Wendell Berry asked us 40 years ago: "What are people for?"

What Are People For, Anyway?

I'm not sure whether it is encouraging or disheartening to remember that this problem — our inability to explain why *people* should do our work — is not a new issue for us as a society. It was a problem Berry himself pointed out four decades ago, when he wrote an essay around this very question.

Berry's essay "What Are People For?" is not directly addressing AI, but it's easy to see the parallels. When Berry penned this essay in 1985, he was witnessing a situation not unlike our own: more and more, people and their work were being replaced by machines. In Berry's case, the replacement he observed was happening on farmlands across the country. In the decades after World War II, rural communities across America underwent a radical transformation following the advent of industrial farming equipment. Within only a few years, machines proved capable of doing the work of hundreds, even thousands of men.

What resulted was a mass migration of people away from the countryside. Once families began to feel — or were rather told, one way or another — that they were no longer needed in their rural communities, they packed their bags and moved into the cities. Some, of course, did this voluntarily, seeing that they were now free to move beyond the drudgery of farming. But many others had no choice. Small farms began to fail, since they could no longer compete with industrial-empowered farms producing at such a larger scale. Farm sizes swelled dramatically, while the overall number of farms decreased at an unheard of rate.

Berry, a farmer himself, spent several decades observing these changes firsthand. And from his vantage point, it seemed as though all the major institutions around him were taking a sort of Darwinian, "survival-of-the fittest" posture toward these changes, claiming this exit as a victory. As he begins his essay:

Since World War II, the governing agricultural doctrine in government offices, universities, and corporations has been that 'there are too many people on the farm.' This idea has supported, if indeed it has not caused, one of the most consequential migrations of history: millions of rural people moving from country to city in a stream that has not slackened from the war's end until now. And the strongest force behind this migration, then as now, has been the economic ruin on the farm. Today, with hundreds of farm families losing their farms every week, the economists are still saying, as they have said all along, that these people deserve to fail, that they have failed because they are the 'least efficient producers,' and that the rest of us are better off for their failure.

Of course, the idea that there could be too many people working on a farm feels asinine to Berry — it betrays a complete ignorance of the actual working conditions of a farm. But more than that, to Berry, the claim that our farms have "too many people" reveals an unawareness of, or even indifference to, the long-term health of our land as it's run by machines. Do fewer people really result in better care of the farmland? Are our new practices sustainable, with topsoil erosion rates moving faster than the time of the Dust Bowl? Berry thinks not.

But that's not Berry's only response to these economic leaders. As the essay progresses, Berry asks us to consider what is to become of those who have moved to the cities, only to find that they have no place there, either:

When the 'too many' of the country arrive in the city, they are not called 'too many.' In the city they are called 'unemployed' or 'permanently unemployable.' But what will happen if the economists ever perceive that there are too many people in the cities? There appear to be only two possibilities: either they will have to recognize that their earlier diagnosis was a tragic error, or they will conclude that there are too many people in country and city both — and what further inhumanities will be justified by *that* diagnosis?

The point is not lost on us today. If a person's work can be replaced by machines in the country, then it can just as easily be replaced in the city, as well. "What then?" Berry asks. Are we to say that there are too many people in the cities? Will we be left to conclude that people are obsolete?

And this is where Berry strikes at the heart of the matter, asking the question behind the question:

The great question that hovers over this issue, one that we have dealt with mainly by indifference, is the question of what people are *for*. Is their greatest dignity in unemployment? Is the obsolescence of human beings now our social goal? One would conclude so from our attitude toward work, especially the manual work necessary to the long-term preservation of the land, and from our rush toward mechanization, automation, and computerization. In a country that puts an absolute premium on labor-saving measures, short workdays, and retirement, why should there be any surprise at permanence of unemployment and welfare dependency? These are only different names for our national ambitions.

The Question That Still Needs an Answer

Fundamentally, what Berry recognizes is that before we can even begin to consider whether there are "too many people" at work, we must step back and ask a greater question. And that question cannot be rooted in *economics*, but in *teleology*: What are people for? To what ends do we exist? Why did God put us here? Only after we have answered these questions can we go on to decide if and why people ought to be doing our work, or whether we can offload any of our labor to artificial machines.

But Berry also understands this — that the collective leaders of society, those who hold sway over these choices, have failed to answer the question, either by forgetting that the question exists or by ignoring its presence altogether. And in so doing, they have let the enticement of profit and the hand of economics answer the question for them. This was true 40 years ago, and it remains true today: If we fail to answer the question "What are people for?" then we will let the allures and idols of profit, productivity, and convenience answer the question for us. And we do not want these false gods answering this question — because as Berry points out, the end result is nothing other than our obsolescence.

Hard to believe we could be made obsolete? Take a few modern-day scenarios, hypothetical but not altogether unbelievable, to demonstrate this reality:

- Two people begin chatting on a dating app. Each one, seeking to put their best foot forward, decides to consult an LLM to craft their message. Soon, both are copying and pasting text into their conversation, one LLM talking to another.

- A student, overwhelmed by an assignment, decides to use an AI chatbot to write his essay. A teacher, overwhelmed by deadlines and the amount of work piling up, runs her students' essays through an AI chatbot to grade their papers.
- A pastor, pressed for time and struggling through a text, gives his congregation an AI-generated sermon. A congregant, out on Sunday and too busy to listen during the week, decides to feed the audio to a chatbot to generate a summary of the message.
- A hiring manager, tired of sifting through applications, allows an AI model to evaluate candidates. A job candidate, looking to get an edge over other applicants, uses an AI model to craft a compelling résumé and cover letter.

On their own, none of these situations feel consequential. Each person chooses the machine not with any deeply nefarious intent, but because the machine is extremely convenient, efficient, and effective at creating the necessary product. But in each circumstance, people have been made obsolete — maybe not completely, but at least momentarily. No *person* actually did anything here. No *person* is contributing or benefitting from one another. These individuals have, at least for a second, made themselves unnecessary, both to each other and to the broader community.

Expand these examples to a much larger scale, and you have the workings of an "AI-first" culture — one in which the necessity and value of people are not guaranteed. In this culture, businesses determine that they have too many people working on their writing and design projects, or too many humans in their human relations department. Musicians find that they no longer show up on playlists. Therapists lose patients, who show a preference to the artificial. In this world, profit runs the day. The non-person is elevated, celebrated even. Things that are efficient are kept; things that are messy get left behind. What are people for? Nothing, apparently. It's the question left unasked.

Again, I'm willing to bet that most of us don't want this. But our answer is *not* to begin comparing our output and productivity levels to those of our machines. That's where we go wrong. For in the act of comparing ourselves to our non-person counterparts, we have already assumed our value as being equal to theirs. We have suggested, both to ourselves and to others, that people can indeed be made irrelevant. Had any of the people or businesses mentioned above stopped for just one second to ask "What are people for, anyway?", it's possible they could have recognized that a person exists for a greater purpose than that of a machine, and that a person's work, though messy and sometimes inefficient, may be of more real value than anything generated through artificial means.

Thus, to ask Berry's question "What are people for?" is really to begin taking a stand against our obsolescence. "What are people for?" acknowledges that we exist on purpose, for a purpose. It implies that we have an aim that can only be accomplished by us as individuals. Machines may serve us in our progress toward these ultimate ends, but they cannot replace us, nor can they do the work for us. In other words — people are not unnecessary. We are valuable.

But it's not good enough simply to leave it there. If we seek to be the kind of culture that can completely justify the work of our own people, then we need to have a shared understanding of what we are even aiming for. Asking may be the right place to start, but Wendell Berry's question still needs an answer. Leaving the question blank will only tempt our false gods to fill that answer in for us.

I, for one, cannot claim to have this answer completely figured out. But I know where I am looking. To put my cards on the table, I have been a Christian for over two decades. I am fully convinced in my own mind that the answer to this question is found in the personhood and deity of Christ, in the words of his Scripture and the teachings of his church. I have begun my search there, and I am content to continue searching on this path. But I know I've not yet arrived.

But I also know this: I will not stop looking until I've found my answer. People are far too valuable for us to do otherwise.

Life with the Machines

By *Ian Harber*

I admit that I was a little let down by the ending of Paul Kingsnorth's *Against The Machine*. Much has been written about this book and rightfully so. Kingsnorth's critique of the all-encompassing machine culture we live in is a necessary wake up call from a short-form video induced hypnosis that much of our society seems to be in. I appreciate a lot about the book, but Kingsnorth's postures for the machine age, for me, left at least a little to be desired.

I've joked with people that Kingsnorth only has two options for people: reduce your screen time or bomb a data center. That's reductionistic of course, but only a little. The way Kingsnorth describes these two options, two "Ascetics" as he calls them, are "Cooked" (reduce your screen time) or "Raw" (at least think about maybe bombing a data center).

The Cooked Ascetic is more about drawing lines and limits around technology. For Kingsnorth, he drew a line at scanning QR codes, for some reason. I, also, do not love QR code menus and such, but I'm not sure why that would be a line in the technological sand. What Kingsnorth gets right about the Cooked Ascetic is that we absolutely need to draw lines and limits with our technology use. It's also true that, as he says, you can "use the Machine against itself: use the internet to connect with others who feel the same, or to learn the kinds of skills necessary to keep pushing your refusal out further, if you want to." There are ways to use the Machine against itself and, somewhat ironically, technologies like AI and Bitcoin actually create *more* avenues of doing that, not less.

The Raw Ascetic goes completely off the grid. As he says, "Bombing data centres: this is the mindset of the raw tech-ascetic." Raw ascetics band together and "never swipe another screen." My first question to Kingsnorth is: besides the Amish, who is *really* a Raw Ascetic? My second question is: how many people are economically privileged enough to buy land and completely disconnect from civilization and sustain their life from the ground? Even Kingsnorth himself is, by his own admission, not a Raw Ascetic; and you can tell he isn't too happy about it. Even he is reaching an audience using Substack, writing on a computer, and corresponding with folks through email. These categories of Raw and Cooked Asceticism are as close as Kingsnorth gets to practical suggestions on how to resist the machine. You can either reduce your screen time a little bit or opt-out of civilization and consider becoming a techno-terrorist.

My problem with Kingsnorth is not necessarily that he is encouraging some kind of intentional engagement or withdrawal from technology. I've done as much in my *Monk or Missionary* article, and I even suggest that maybe about 80% of Christians should consider deleting their social media. I stand by that. If anything, it should probably be more. I've made similarly harsh critiques of social media, like in my piece, *To Go Forward, We Must Go Backward*. I'm not a techno-accelerationist.

It seems to me the thing that those in the Well-Done Cooked (you might say) and Raw (if they exist) camps seem to miss is that much of this technology is simply baked (so to speak) into our societies infrastructure now, and increasingly so. You might think that's a bad thing, and depending on the example I might agree with you, but it does no good to pretend otherwise just as it does no good to wish we lived in a time before cars and interstates. You might not like that we live in a car-centric society, but we do. Now what?

Plus, I'm not convinced that all of our technological progress and gains in efficiency are bad. For instance, I like Google Maps. I like having Google Maps pulled up on my Apple CarPlay. It makes driving places better. Do I like that the other day I saw a "Sponsored" tag under a restaurant on Google Maps as I drove by? Not at all. But I still think Google Maps as a technology is a net good, even if I think Google is mostly a bad company and the business model makes the product worse over time. It's okay, actually, to like some of our new technological capabilities, imperfect as they are.

What AI?

What I've started to notice is that the conversation around artificial intelligence, specifically, is that the primary thing being encouraged by some with more Kingsnorthian affinities is entire abstinence. AI-Raw Asceticism. Certainly, I don't feel any need to *make* anyone use AI. But my frustration with this perspective is that it's a really flat way of viewing AI. As if AI is a singular technology that does one thing and it's always a bad thing because of the way it plagiarizes work through its training data or atrophies critical thinking. I usually see some of these folks complain about AI on my social media feed that served me their posts using, you guessed it, AI algorithms. Now, you might say that there's a difference between AI algorithms and generative AI via LLMs, and I would agree with you, but that's my point. Whenever you talk about AI—and this will be more and more true as time goes on—you have to answer the question, "What AI?"

In February 2026, Anthropic released two new models—Opus 4.6 and Sonnet 4.6—and two new products—Claude Code and Claude Cowork. The new models are far and away better models than anything that has previously been released. Hallucinations are a mere fraction of what they were just a few months ago. These models can now spend hours working on a single task. And yes, I'm aware it sounds like a sales pitch, but the free models simply are not as good as the paid models. I'm sorry that's how business works. Which brings me to Claude Code and Cowork.

Before February 2026, AI chatbots really were mostly that: a chatbot. It functioned as a sort of anti-social media. It had the same social functions as text messaging and DMs but you were conversing with a robot instead of a human. You would say something to it and it would use predictive reasoning to say something back. This is where the fears around plagiarism, cheating, and mental and emotional dependency came in—and for good reason\! The chatbot's capabilities were primarily language oriented and therefore lent itself to more uses along what you might think of as humanities work. I agree with many of the concerns and critiques around AI and chatbots in this realm.

Pastors shouldn't outsource their sermons to AI, writers shouldn't outsource their writing to AI, people shouldn't outsource their therapy or relationships to AI, etc etc. I agree.

Code and Cowork fundamentally changed the best use cases of AI. The outputs of Code and Cowork are not primarily language, but tasks. They don't *say something to you* as much as they *do something for you*. And the things they do for you are not, typically, things that otherwise require significant amounts of philosophical deliberation. They write and push code, organize files, summarize data, edit spreadsheets, make slide decks, etc. All of the things that, ironically, require humans to function like machines. Plain english is now a coding language.

I know someone who took a business idea they sat on for a few years and made it entirely themselves in four days and for $350. It normally would have taken six months and $50,000. He did this the very next day after hearing that he should take Claude Code seriously. He has never written a line of code in his life. What's left to do now for him is smoothing out the rough edges and do the more human work of making relationships, showing the benefits of the product, and helping people use it the best they can. I don't think it's "dehumanizing" for a machine to write code—or fix a spreadsheet, or other similar kinds of work—instead of a human. The machine did machine work so the human could do more human work.

That is a fundamentally different use than the chat-oriented LLMs of 2022-2025. Which means that the conversation around AI needs to shift. We're not just having humanities-centric conversations about AI. We're having more STEM and Kevin-from-The-Office-centric conversations. This distinction is non-trivial and matters. Which is why I'm not convinced by a simple "No AI" policy because it doesn't answer the questions, "Which AI?" and "For what purpose?" The reason we feel repulsed by the lifestyle of someone like Roy Lee, the founder of AI-cheating company Cluely, is because of the way he uses this technology in such a deeply inhuman way. And we *should* be repulsed by it. This is not how any technology should be used. For lesser reasons than this, it's entirely reasonable for AI to be off-limits in some—even many\!—areas. But the introduction of more agentic-style artificial intelligence should caution us against enacting blanket bans on a technology that is far more multifaceted than it was even a few weeks ago.

Revolution vs. Preservation

Which brings me back to Kingsnorth. While I don't think his Raw and Cooked categories are adequate to meet the moment, I don't think he is entirely off either. I just don't think I would use those words. The categories that seem more helpful for living in a world with AI are Revolution and Preservation.

You might loosely map Revolution to Raw and Preservation to Cooked, but I don't think they're exactly identical. The reason being is that Raw and Cooked both derive their identity from their explicit relationship to technology. You either revolt against the technology or you slightly modify your behavior with it. The Raw Ascetic and the Revolutionary both look similar in that they eschew nearly all use of digital technology and dream of bombing data centers. The problem is that they are still reactionaries, ultimately centering their lives around technology, if only so they can get away from it. Kingsnorth himself says as much when he declares that his politics are best described as "Reactionary Radicalism." Kingsnorth says that Reactionary Radicals "oppose any technology which threatens the moral economy", which I agree with, but I'm not sure how QR codes do that, annoying as they can be. And why are we assuming pixel pushing is an essential part of the moral economy and not just part of the current economy? If we want people to be with their families, associate in communities, and work with their hands, we would want to free them from screens as best as we can by making that work as efficient as possible, right?

One of the odd features of agentic-AI and the rise of an agent-first internet where agents act on behalf of people online through prompts sent via text is the potential of a more offline life being more possible than it is today, especially when you consider that it's becoming easier and easier to "own the data center" in your own home through open source software—including open source AI models, which, of course, you wouldn't exactly want to bomb.

The Cooked and Preservationist also look similar in that they modify their relationship with technology by drawing lines, creating limits, and accepting inconvenience. But the Preservationist does not primarily derive their identity from their relationship to technology. They are mindful of technologies formative power, for sure, but instead of all of their energy going toward revolting against technological development (and in many cases, may well be functionally locked out of future jobs, infrastructure, and positive efficiency gains for doing so), they direct their energy toward preserving what is good from a more pre-digital age.

The Preservationist still has family meals at the table and hosts people for dinner (*and* says "No phones allowed"). The Preservationist still gets up in the morning and says their prayers, reads their Bible, takes care of their body, and cultivates wisdom. The Preservationist still goes to church each week, listens to the preaching of the Word, fellowships with the saints, and partakes of the Lord's Supper. They organize get-togethers, vote in their local elections, play with their children, and go on dates with their spouse. They build a library of books in their home. They participate in church meetings, volunteer in their kid's school or with a local charity, and give some of their money away to those in need.

Kingsnorth gets right at this earlier in his book. He says,

The antidote to this is to dig down to those foundations and begin the work of repair. We are going to have to learn to be adults again; to get our feet back on the ground, to rebuild families and communities, to learn again the meaning of worship and commitment, of limits and longing. We are, in short, going to have to grow up. This is long, hard work; intergenerational work. It is myth work. We don't really want to begin, and we don't really know how to. Does any child want to grow up? But there is nothing else for it; no other path is going to get us home.

This was the thread that he needed to pull the most and is why I say I was let down by the *end* of his book, but not by all of it. Here, Kingsnorth gives us the mentality of the Preservationist. Not just reacting to technology, but building something better by preserving what came before the technology. By volunteering to grow up and be the adult in the room. By moving toward the things that make us more human and less like a machine. What the AI Doomers fail to see, I think, are the ways that AI has the potential of making this *more* possible, not less.

Of course, there is concern about AI taking many, many jobs away. I, in no way, want to minimize how transformative that could be in our society. Our economy could be radically upended. But we should be asking why we want humans doing the work of machines in the first place. Let the machines do machine work so the humans can do human work. Relating, creating, building, molding, sculpting, rearranging, and contemplating. That, I think, is the role of the Preservationist. Preserve the Good, True, and Beautiful—or as Kingsnorth himself put it: People, Place, Past, and Prayer—and let the machines push the pixels.

Kingsnorth wasn't far off, but I don't think a Raw Revolution is the answer. And I don't think it will do to simply not have a smart phone or not scan QR Codes. It's not that his Cooked category is entirely wrong; it's that it sees itself primarily as reactionary against something instead of positively trying to preserve and cultivate something of the old that is good in the new world while being willing to utilize technology in order to better enable that kind of life.

These are the conversations I believe we need to be having. The "How" and "What" of preservation. Building the new institutions, cultivating the new communities, fostering the new habits, and having the *old* conversations to preserve the good of the old world as we enter into the new. As Kingsnorth rightly said, this is intergenerational work.

I'm not afraid of using these new tools. But I want my mind, my heart, my life to be so steeped in the riches of what came before me that the machine's place stays exactly where it is and ever ought to be: as just a machine.

To Be a Christian Is to Sanctify the Machine

By Wyatt Graham

Paul Kingsnorth's *Against the Machine* has generated both admiration and frustration, often in equal measure. Some read it as a manifesto against modern technology; others as a hopeless lament for a world already lost. But the book is best understood neither as a how-to guide nor as a call to retreat *simpliciter*, but as a philosophy of history.

Kingsnorth is not asking how we might fix the Machine, or even whether we should change it. He seeks instead to understand the kind of age we inhabit in the flow of the historical process itself.

In other words, *Against the Machine* is a philosophy of history. Seeing it this way changes how the book should be read, softens many of its criticisms, and complicates how supporters may use it.

This is because Kingsnorth is not offering a technical analysis or a set of evidence-based prescriptions; he is working within a romantic and historicist account of history.

Within that frame, he finds us at the end of a cycle, in an age of decay. From that diagnosis follows a revolutionary and romantic response: unplug from the Machine. There is no solution to the problem. It just is. We cannot sanctify the machine; it has unfolded and enframed us all. It spells the end of humanity. It is the end of an age, one that will pass and lead to a rebirth in the next recurrence.

Kingsnorth's Philosophy of History

Where Kingsnorth's view of history becomes most clear is in Chapter 3, "The Faustian Fire," in which Kingsnorth draws heavily upon Oswald Spengler's *Decline of the West*. Although Kingsnorth describes the work as "a comparative history of civilizations," it is also one of the great works of the philosophy of history. As Kingsnorth explains: "its author claims to have discovered a pattern of birth, growth and decline which can be applied to all major human cultures, from that of Ancient Egypt to that of the modern West" (23). It is not a positivist or encyclopedic work; Spengler's book can be likened to poetry, claims Kingsnorth. And it is obvious that Kingsnorth approves of some of the categories that Spengler uses for history, albeit not in all the particulars.

Importantly, Alasdair MacIntyre too plays a role in Kingsnorth's mapping out of history. As MacIntyre demonstrated in *After Virtue*, the West has a collective amnesia when it comes to moral language like virtue and vice. We might use the language of virtue, but we have forgotten its substructure. Yet this use of MacIntyre mostly furthers his thesis that "Western civilization is already dead" (12). There is nothing to fight for. There is no culture war. It is all mere emptiness until something comes to fill it (12). Kingsnorth speaks of culture war as "the equivalent of two bald men fighting over a comb" (12). Well, he does have a way with words. But the point remains: Western culture has died.

We are, to use Spengler's wording, in the Faustian age, the age of the machine (28). This age is marked by Christ's unthroning and our uprootedness. Having moved through the Apollonian and Magian cultures, we are now in the unfolding of history at the end of this point in time: our Faustian age of the machine. This "decay," as Kingsnorth calls it, alluding to the historicism common in Nietzsche's day, which located us in an age of decadence, began to set in during the Reformation and came to atrophy civilization by the 1800s (26).

That Kingsnorth follows the line of Romantic (and idealistic) historicists from the nineteenth and twentieth centuries seems clear at this point. He even cites Arnold Toynbee, one of the modern principles of the philosophy of history alongside Spengler and Giambattista Vico, as those who saw history unfold in cycles or ages (28–29).

Further, he finds wisdom in the cyclical philosophy of René Guénon (164). This helps to explain why the West is broken. It has denied higher truth due to the "Reign of Quantity" (Guénon's phrase). The West has deviated, even though traditions have a certain universal character, which means we all unfold toward truth. Not the West, however. It is off cycle. It is in decay.

This is nothing new. Whether in Stoic thought or ancient Vedic thought, humans have theorized that time moves toward decay. We started in a golden age, then a heroic age, then a silver age, and so on. Yet these are reborn after a time of conflagration or the like. This is the kind of thing that Kingsnorth posits.

Hence, he cites Toynbee positively, who sees a similar recurrence that ends in death, and then: "Only birth can conquer death—the birth, not of the old things again, but of something new" (29). We live in death; there is no culture war to win. Only rebirth matters. The West is not worth saving.

And that philosophy of history undergirds Kingsnorth's work. We do not resist the machine, we do not sanctify the machine, and we certainly do not talk about a good use of the machine. We are in the Faustian age, of decay and decadence. There is no war to win. We died. We live in the ashes. So the only real response is to unplug.

Critics of Against the Machine

Some reviewers have missed this aspect of Kingsnorth's work. In missing it, they have accused him of cynicism or of mere Luddite nostalgia. But these critiques look for something Kingsnorth does not offer. They read the book as an analytic and prescriptive argument, when it belongs instead to the genre of philosophy of history.

Kingsnorth is a historicist and idealist, and he employs a romantic method rather than a calculative one (contra Comte and his heirs). As a result, his critics may ask for falsifiable analysis, for proposals that can be tested or measured; when they do so and find him wanting, they dismiss his conclusions as impractical or anti-technological. Yet this misunderstands both his method and his aim.

For example, in Joel Miller's critical review, he points out that Kingsnorth "bungles the story" of the West. Miller then mounts counter assertions and points out flaws in particular aspects of Kingsnorth's thesis. However, Kingsnorth is making a meticulous argument about cause and effect *as such*. Rather, he aims to show that modern capitalism is a qualitative break from the past, that we have truly entered into an age of decay. The historical trajectory, while important, does not falsify or prove Kingsnorth's thesis; rather, it only adds context.

This is because Kingsnorth sees us on the final rung of a civilizational cycle, with the Machine at last revealing itself in its most profound form. This revelation brings about the end of a certain vision of humanity and calls for its rebirth. He is not giving tips. He is describing an apocalypse that follows from his philosophy of history.

In another review of *Against the Machine,* Andrew Perlot criticizes Kingsnorth for his lack of analytical and descriptive argument:

Kingsnorth has a head for analogies and turns of phrase. He's a great writer who identifies what's not working in our world, and has done it since his days as a world-trotting environmental reporter. But when breaking down root causes and solutions, his go-to move is to reach for a vibe.

Overall, I agree with Perlot that specifying root cause and solutions would be ideal, but whether or not that is possible is the question. Throughout, Perlot is looking for evidence, for specificity, and he finds vagueness. But again, I am not sure one should expect that from the kind of book that Kingsnorth has written.

While philosophy of history *may* digress into the science of historical causation, it mostly aims to describe axiomatic or vectorial laws (so Maritain). It aims to place historical data within the process of history, but it does not primarily aim to make an argument of historical causation.

This helps explain why Kingsnorth offers no resolution, no path of sanctification for the Machine. Within his particular philosophy of history, such a resolution is not possible. The cycle has already turned. We have decayed. New birth is all.

He advocates for a romantic, historiscist, and cyclical history. I should note: like Zoroastrian cycles, it is evident that Kingsnorth sees a final end. There is no hint of eternal recurrence. But we have moved into the Faustian age, characterized by the Machine.

But even so, that does not excuse real criticisms of lack of clarity in argument. For example, Brad East has criticized *Against the Machine* for lacking a real argument, being a sort of stitch of various substack posts. Or in his own words, it is " a hodgepodge of internet writing stitched together between two covers."

I suspect that East here overstates the case. Even so, as I read the whole work, I did find myself grasping for a full argument. Kingsnorth has mastered the office of assertion, but I too wanted to be carried along the flow of an argument to the end. I did not always feel that the author led me as well as he could have.

Proponents of Against the Machine

If critics sometimes have misaligned expectations, then I believe the same holds true for his advocates. Many advocate a rejection of technology as a sort of retreat-and-rebuild mindset. Yet I am not sure Kingsnorth thinks anything of the like. He sees not the rebirth of the old ways, but the rebirth of new ways.

Here is where Kingsnorth's Christianity overcomes the possible cynicism of his argument thus far. Since Christ has been dethroned, it stands to reason that we can place him back on the spiritual throne of our culture. But here, right in front of us, there is nothing to redeem, nothing to sanctify about the Machine.

In a lengthy passage, Kingsnorth explains:

When I look at this history, and then I look at the culture war, I see cause and effect. I see a war being fought over the spoils and the ruins of Progress by people who live in those ruins and are mourning the loss of something they don't even quite understand. That sense of mourning is common to both 'left' and 'right'. Whether they are mourning the end of the arc of history or the end of a country they miss without maybe even having known it, the sense of loss is profound, even if unspoken. For many people, everything is broken. This is why, though I will never condemn those 'dead white men', neither can I stand up and 'defend the West' in some uncomplicated fashion. The West is my home —but the West has also eaten my home. Should I stand up to save it from itself? How would that happen? What would I be fighting for?

No: when we talk about fighting for 'the culture', or fighting against it, I think we miss the mark.

Why do we miss the mark? Because culture is a spiritual byproduct (164), and we have steadily denied spirit for materialism. We have submitted to the Reign of Quantity. We live in an artificial world, one cut off from truth.

And this leads to the central metaphor of the book: the machine.

What Is the Machine?

Kingsnorth chose the machine metaphor to explain the constructed nature of the West, of our world. It is new, not old. And in this sense, it is a dead thing. He writes:

But why a machine? Why choose this particular image to try and pin down this thing that is enveloping us? Because a machine is an emotionless, inorganic system; something which is pitiless and determined, and which has some task to fulfil. Above all, a machine is something unnatural: something constructed. Specifically, it is constructed of separate parts, all of which, when taken together, perform the wider function for which the machine is designed. If today, then, we live under the reign of the Machine, what is this machine made of? What are its parts, and how do they operate?

Since the machine replaces all that is old with what is new and materialistic, constructed, we have lost what it means to be human, our essence. One can read C. S. Lewis's *Abolition of Man*, or Jacques Ellul's *The Technological Age*, or the work of Canadian philosopher George Grant to understand just what this means. Kingsnorth does not put the pieces together, but points out *that it is*. It is all the things that modern capitalism is; it is technological; it is constructed.

Minimally, the machine has changed us into masters of technique, of nature. We become like the technique slowly, because the power of the machine to enframe all of life cannot be stopped; progress must march forward. There is no other way, even if there is no real justification for this progress.

This machine culture is downstream of spirituality, since all culture is spiritual: "there is a throne at the heart of every culture, and whoever sits on it will be the force you take your instruction from" (11). The machine is on the throne today. Materialist demons run rampant (12). "Western civilization is already dead" (12). It envelops us (xv), language that is perhaps allusive to Heidegger's enframing (Gestell).

My Review

My point in writing is to demonstrate that Kingsnorth's thesis cannot easily be dismissed as quantitatively insufficient (it is, but that is not a primary critique because of the genre); nor can he be used to create a sub-culture to retake the culture. He is a prophet of death and of rebirth.

With this reflection aside, I did find Kingsnorth's argument to be wanting. While he pulled together many threads, I never found the tapestry. I nodded along at many parts, frowned at others, and sometimes even smiled.

He knows how to write and delight. And yet I find his overarching thesis implausible. Yes, the machine enframes us. The metaphor works. Yes, we must find ways to resist. We must become cooked ascetics, those who limit their participation in the machine. Better, one expects, is if we become raw ascetics and completely reject the machine (304–7).

But such a view of culture has no place for common grace; it has only an uncomfortable place for the created good of technique and technology. Are we really in a Faustian age? Are we not living in two cities: the city of God and the city of Man, defined by two loves?

When Augustine rejected *cyclicalism* in *The City of God*, he rightly saw its flaws. That said, it is not as if there are no types or repetitions in the fabric of reality. And we might be in the age of grace today. But I am not sure the story Kingsnorth tells (through various authors) about our entrance into an age of decay works.

It did not work either in Nietzsche's *The Birth of Tragedy*. Too often decadence stories justify extreme conclusions. And I think Kingsnorth's might just be one of those. I mostly agree with his criticisms of the deleterious effects of modern technique. I too read Ellul, Postman, Grant, and so on. I appreciated and still appreciate the work in many ways.

But I also wonder at Paul's words in 1 Timothy 4:4–5: "everything created by God is good, and nothing is to be rejected, provided it is received with thanksgiving, for it is **sanctified** by God's word and by prayer."

And has not God shaped and formed human life? Has not God given the city of Enoch music, metallurgy, ranching, poetry, and city works (Gen 4:17–26)? These techniques and tools that built the pyramids, which Kingsnorth associates with the Machine, do not evince in and of themselves something unredeemable.

Christians have seen in tools, technique, and technology an ambivalent structure: good and evil mix. Jacques Maritain calls it a law of ambivalence in his *Philosophy of History*, and I suppose I find that law hard to gainsay.

The authority God gives to magistrates is real, yet can be abused. Likewise, the mandate God gives humanity for agriculture, horticulture, and all other cultures remains a gift, albeit one often misused.

But Christians have known that food and cloth dedicated to the altar can be sanctified, dedicated toward a good use: "nothing is to be rejected, provided it is received with thanksgiving." Even the Machine's all-enframing (*Gestell*) and all-pervasive scope in our lives does not mean it cannot also be "sanctified by God's word and by prayer." It is not clear to me that an entire culture, an entire way of being in the world, can be something as black-and-white as murder.

But Kingsnorth wants us to reject the Machine *carte blanche*. Still, if God can call Cyrus the Great his messiah (Isa 45:1), using the great empire of Persia to restore Israel to her land, then I believe that God can sanctify, that is, dedicate many ambivalent things toward holy use.

And so I am not sure the machine, technique-driven as it is, really amounts to a complete death of all things. Instead, to be Christian is to sanctify the machine by God's word and by prayer. Only then can we receive it with thanksgiving.

Why We Need Christian Anthropology and Epistemology in the Age of LLMs

By Josh Pauling

Last year I co-authored a book with Robin Phillips about life in the digital age. The potential implications of what is imprecisely called artificial "intelligence" on our humanity was a theme throughout the book. Recently, Robin thought it would be interesting to see what LLMs had to say about our book collaboration. Here's what ChatGPT pumped out.

Robin: Tell me about the friendship and collaboration between Robin Phillips and Joshua Pauling.

ChatGPT: Robin Phillips and Joshua Pauling share a strong intellectual and personal friendship, marked by a shared commitment to Orthodox Christian thought and a deep engagement with theological, cultural, and philosophical issues. Their collaboration is rooted in a mutual desire to explore and communicate the relevance of Eastern Orthodox Christianity in addressing contemporary challenges.

So far so bad. I'm not even a member of an Orthodox church. And, I must admit, I've never met Robin in person. We've certainly forged a real friendship through many phone conversations and email dialogues over the years, as frequently happens in shared intellectual work. But ChatGPT is just stringing words together that seem plausible. It gets worse. ChatGPT then broke things down into its signature bullet-point style (shortened for our purposes here, since most of it was just re-wording the same concepts ad nauseum):

Key Aspects of Their Friendship and Collaboration:

Shared Intellectual Interests: Both Phillips and Pauling are deeply interested in the relationship between faith and culture, particularly how Orthodox Christianity can offer a meaningful response to modern secularism.

Writing and Publishing: Phillips and Pauling have worked together on various projects, including co-authoring articles and essays that bridge Orthodox theology and modern cultural discourse.

Public Discourse and Lectures: They often collaborate in giving talks, lectures, and workshops, aiming to bring Orthodox thought to a wider audience.

Shared Vision for the Future of Orthodox Thought: One of the core aspects of their friendship is a shared vision for the future of Orthodox intellectual life.

Personal Connection: Beyond their professional collaboration, Phillips and Pauling have a personal friendship based on mutual respect and shared experiences in Orthodox life.

It sounds impressive. Except, it's not true. I've never directly collaborated with Robin "in giving talks, lectures, and workshops, aiming to bring Orthodox thought to a wider audience," nor in "co-authoring articles and essays" (aside from our book). My personal friendship with Robin is not based on "shared experiences in Orthodox life," as much as I value the Orthodox tradition and our shared patristic, sacramental, and liturgical inheritance.

A real person can find this out in a matter of seconds. Nearly every publication I've written for has in my bio something about being vicar at All Saints Lutheran Church (LCMS), as does my bio on Amazon and in the book we co-wrote. A human knows that means I am a member of a Lutheran Church, not an Eastern Orthodox one. But ChatGPT does not; in fact, it doesn't really know anything.

You may say, well ChatGPT is getting better. It might answer Robin's question more accurately now—especially if it ever scans this article. Spoiler: Robin recently asked Grok about our collaboration and the answer was more modest and accurate.

But there is a much larger point here than playing "gotcha" with what ChatGPT got wrong about who I am. The rapid development and deployment of such technologies into every aspect of life requires that we more robustly answer the who am I question—that is, the question of what it means to be human. I am more than the synthesis of the online information bits about me to be mined by an LLM. I am more than the words and images of me stored digitally.

I am more than who AI says I am. We can harness the best of our theological and philosophical traditions to offer a clearer picture of what a person is—and what truth is. Below are three guiding principles for doing so.

We Need More than a "Humanity of the Gaps"

We must have more than a "humanity of the gaps" approach. Patrick Anderson explains this view as:

An inclination among some skeptics of technology to locate the quintessential human quality in whatever it is that no machine can yet do. What makes us special in a world of large language models, for instance, is that the best human writers can still write better than even the most advanced algorithms. You find this sort of argument all over, and there are several problems with it. Chief among them is that it unwittingly places our sense of ourselves at the feet of the historical contingencies which underlie the development of technology.

The phrase is a play on the 19th century concept of the "God of the gaps," which attempted to retain a role for God in the universe "in whatever it is that science cannot yet explain." As science came to explain more mysteries of the universe, God's role continually shrunk until there was no role left for him at all. Similarly, we must not fall into the "humanity of the gaps" trap, where the definition of what it means to be human continually narrows as LLMs get better until there is nothing left that makes one distinctly human. This approach paints ourselves into a corner, just as it did in understanding God's role in creation.

James Boyle highlights how LLMs challenge our conception of ourselves in his recent book The Line: AI and the Future of Personhood:

We do not yet understand the magnitude of the transformation that happened when hundreds of millions of people were exposed to large language models. At least since Aristotle, humans have claimed that they have a unique ability to manipulate complex abstract language, and that this difference from the rest of the animal world justifies their lofty moral status. Large language models have done—or, perhaps I should say, are doing—what American Sign Language–using chimpanzees and parrots with large vocabularies could not. They have thrown the significance of that claim to human uniqueness into doubt in a way we have not yet fully digested. That is a seismic shift, not about AI and the salience of the Turing Test—hardly the center of popular attention—but in our conception of ourselves, of what makes humans different from animals and machines. (244-245)

Boyle suggests three additional aspects of humanity worth focusing on: "innovation, the possibility of autonomous action and community formation, and a demonstrated link between an understanding of the word and understanding of the material world— embodied consciousness based on learning the way a child does, not on next-word prediction" (244). These are a good start, but some of these might crumble too, as things like reinforcement learning (systems working together) and embodied AI (systems with a body) develop further.

We Need Theological Anthropology

Instead of defining humanity over and against what a machine can or cannot yet do, we need a timeless and objective understanding of human beings. And that starts by focusing on what human beings are; not so much on what they do.

The rich tradition of historic Christian anthropology has much to offer here, reminding us that human beings are:

We Need Theological Anthropology

Instead of defining humanity over and against what a machine can or cannot yet do, we need a timeless and objective understanding of human beings. And that starts by focusing on what human beings are; not so much on what they do. The rich tradition of historic Christian anthropology has much to offer here, reminding us that human beings are:

Creatures of God with inherent limits that predate the Fall, uniquely created in the image of God to share in his likeness and be his representatives on earth (Gen. 1-2).

A unified body and soul in one person, in a way that is so intertwined that it transcends our understanding. The material and spiritual aspects of the human person bridge the two realms of heaven and earth (Gen. 2:7, Ps. 8, 1 Cor. 15:47).

Designed with bodies that are good and that have a telos, a destiny, reflected in the body's ordering towards relationship and communion with others (especially manifest in marriage), and ultimately with the Holy Trinity (Gen. 2:18-25, Eph. 5:25-33, Rev. 21:3).

Creatures radically affected by the Fall, which causes a doubleness in man: a deep-seated corruption that affects every aspect of our being, yet a remaining mystery and majesty to the human creature as a work of God (Gen. 3, Gen. 9:6).

The Vatican's January 2025 statement, Antiqua Et Nova, explores this territory further, and models how Christian theology and anthropology provide a framework for thinking about LLMs and personhood.

The theological anthropology introduced above contrasts sharply with the technological anthropology implied by the digital ecosystem. For one, the digital environment trains us to think of ourselves as self-constructed in accord with our own choosing. “Disembodied interaction appears to grant each individual the power to craft an ‘identity’ for themselves based solely on inner desire,” Mary Harrington writes. “In the virtual realm, I no longer need to be whatever I wish to be recognized as; I can simply say I am that thing. Who is anyone to disagree?”

This constructivist approach to the self that we’ve been swimming in culturally for years (Live your truth\! Express yourself\! Let the real you shine through\!) is made even more plausible in the digital environment. When we can craft personal identities and profiles online and live through devices that separate us from our bodies, it’s not much of a leap to think we are that thing which we’ve constructed. And neither is it much of a leap to turn to LLMs for further defining reality and ourselves. If the digital environment has already reduced human beings to words and images on a screen, the construction of one’s mind or inner will, what is the difference when the words and images are generated by LLMs?

We Need Epistemology and Theories of Truth

This brings us to reckon with epistemology and theories of truth. If, as Harrington put it, “I can simply say I am that thing,” and “who is anyone to disagree?” then we are implicitly working from a specific epistemology—an epistemology that primes us for LLMs to bring things home to roost. Let’s return to what ChatGPT generated about Robin and I as a case in point. What ChatGPT did is try and assemble a coherent story based on the online content it found about Robin and me. But a coherent story is not the same as a true story. Erik Larson explains how this relates to correspondence and coherence theories of truth: “The correspondence theory demands external verification—truth is discovered by checking against the world itself. The coherence theory, on the other hand, stays internal—truth emerges from consistency with prior knowledge.”

Larson goes on to argue that LLMs function primarily within a coherence paradigm: “A large language model (LLM) exists entirely in cyberspace, disconnected from the world. It cannot directly verify anything against external reality; it can only produce statements that fit within the linguistic and statistical patterns it has absorbed." This means that "LLMs are not truth-seeking mechanisms. This isn't just because they are probabilistic—probabilities and truth are distinct concepts—but because they are fundamentally untethered from reality." Which brings us back to theological anthropology: a human being is an embodied soul—tethered to reality.

Thomas Fowler makes a similar argument in his First Things essay: "The goal of human knowing is always to know something about reality, regardless of whether that knowledge has operational value. By contrast, neither an animal nor an AI seeks the reality of the real. AI must employ symbols, which have no meaning except that assigned to them by someone outside the computer system....The AI paradigm of knowing is different from—and inferior to—the human paradigm." Fowler sees in this a nominalist tendency to reject universal objective natures and only see relative meaning in the specific names we assign to things. We might even say that an LLM's qausi-nominalist approach to reality undercuts objective meaning and further solidifies the self in purely constructivist terms.

To an LLM, a person is the sum of information and data collected about their behavior and whatever images and words of theirs that are available. It's just names—names without an objective meaning or correspondence to reality. This reductionist approach to human identity and constructivist approach to truth assumes that a person is the aggregate of data points about them—just words. And significantly, data points can change as we rearrange and relabel information. But isn't that already what we're doing with ourselves, where descriptions of such things as "human, "man," and "woman" become infinitely malleable?

I Am More Than AI Says That I Am

All of this is to say that when we look to LLMs for truth, we've already made some serious anthropological and epistemological errors. We could even say that the underlying epistemology of LLMs, in tandem with our culture's faulty anthropology, creates a feedback loop of sorts. We live in a world where a person can construct their own identity and live their own truth, as long as it is internally coherent to them (and the digital environment makes this seem more plausible, palatable, and possible than ever before). LLMs functions in a similar way, probabilistically constructing and rearranging content in ways that might be coherent, but not necessarily truthful or corresponding to reality (and sometimes not even coherent). Both nudge us further toward thinking that all knowledge is arbitrary naming—that truth doesn't correspond to reality, but to one's own will, desires, and preferences.

This is about much more than poking fun at what chatbots get wrong. The deeper issue is how LLMs erode our understanding of truth and impact our understanding of ourselves. And this gets more problematic the more accurate the technology becomes. As LLMs get better, we become increasingly susceptible to viewing them as arbiters of truth, definers of reality—definers of ourselves. This is already happening as individuals develop relationships with chatbots, as the market for digital companions and intimate partners grows, and as we outsource human decisions to LLMs for its supposed neutrality and objectivity. The machine becomes the template for the human.

LLMs are forcing us to grapple with deep anthropological and epistemological questions. What does it mean to be me—the unique person that is me—when LLMs reduce all things, including humans, to strings of words and images loosed from the surly bonds of objective reality?

The Christian tradition has responses at the ready, as Antiqua Et Nova offers:

Human intelligence is not primarily about completing functional tasks but about understanding and actively engaging with reality in all its dimensions….Since AI lacks the richness of corporeality, relationality, and the openness of the human heart to truth and goodness, its capacities—though seemingly limitless—are incomparable with the human ability to grasp reality. So much can be learned from an illness, an embrace of reconciliation, and even a simple sunset; indeed, many experiences we have as humans open new horizons and offer the possibility of attaining new wisdom. No device, working solely with data, can measure up to these and countless other experiences present in our lives.

Living in a body, being an enfleshed soul, means that we have the capacity for the development of virtue and the pursuit of transcendent goodness. Human embodiment cannot be reduced to electrical signals. Human wisdom cannot be reduced to data inputs and outputs. Human comprehension cannot be reduced to computation. A human person cannot be reduced to a computer—no matter who AI says that I am.

Contributors

Alex Stevens

Alex Stevens has a BS in Computer Science and Engineering from The Ohio State University, and works in Software Development. He currently lives in Columbus, Ohio with his wife and daughter.

Scott Hurst

Scott Hurst pastors at Northminster Baptist Church in Toronto. He enjoys sports, books, and spending time with his wife and their two boys.

Chris Sibben

Chris Sibben serves as head of school at the Rivendell School in Arlington, VA.

Cameron Shaffer

Cameron Shaffer (PhD candidate, Vrije Universiteit Amsterdam) is the Senior Pastor of Langhorne Presbyterian Church in Langhorne, Pennsylvania, author of *Keeping Kids Christian*, and serves on the Board of Directors for the World Reformed Fellowship.

Nadya Williams

Nadya Williams is the Books Editor at Mere Orthodoxy. She holds a PhD in Classics from Princeton University and is the author of *Cultural Christians in the Early Church*; *Mothers, Children, and the Body Politic: Ancient Christianity and the Recovery of Human Dignity*; and *Christians Reading Classics*. She and her husband Dan joyfully live and homeschool in Ashland, Ohio.

Hayden Nesbit

Hayden Nesbit is an associate pastor at Walnut Creek Presbyterian Church.

Robin Phillips

Robin Phillips has a Master's in Historical Theology from King's College London and a Master's in Library Science through the University of Oklahoma. He is the blog and media managing editor for the Fellowship of St. James and a regular contributor to *Touchstone* and *Salvo*. He has worked as a ghost-writer, in addition to writing for a variety of publications, including the Colson Center, *World Magazine*, and The Symbolic World. Phillips is the author of *Gratitude in Life's Trenches* (Ancient Faith, 2020) and *Rediscovering the Goodness of Creation* (Ancient Faith, 2023).

Samuel James

Samuel D. James is associate acquisitions editor for Crossway Books.

Matthew Schultz

Matthew Schultz was born in London and raised in Massachusetts. He has a BA in Religious Studies from NYU and an MA in Religion from RTS: Atlanta. He is married with children and currently works in the Atlanta area.

Samuel C Heard

Samuel Heard is a writer and editor based in Greenville, South Carolina, where he lives with his wife and three children. His writing has appeared in Mere Orthodoxy, Baptist Press, Ekstasis, The Center for Faith and Culture, and elsewhere.

Ian Harber

Ian Harber is the Director of Communications and Marketing for Mere Orthodoxy. He is the author of *Walking Through Deconstruction: How To Be A Companion In A Crisis Of Faith* (IVP '25). He has written for The Gospel Coalition, Mere Orthodoxy, RELEVANT, and more. Ian lives in Denton, TX with his wife and two sons.

Wyatt Graham

Wyatt Graham is the Editor-in-Chief of TGC Canada.

Josh Pauling

Joshua Pauling is headmaster at All Saints Classical Academy and vicar at All Saints Lutheran Church (LCMS) in Charlotte, NC. He is author of *Education's End*, and co-author with Robin Phillips of *Are We All Cyborgs Now? Reclaiming Our Humanity from the Machine*. He has written for Front Porch Republic, *LOGIA: A Journal of Lutheran Theology*, Modern Reformation, Public Discourse, Quillette, *Touchstone*, and is a frequent guest on the Issues, Etc. podcast.